Evaluation Protocol, Not Representation or Topology, Drives Reported EEG Classification of Alzheimer's Disease and Frontotemporal Dementia

Devguru
Indian Institute of Information Technology, Nagpur
bt23csd060@iiitn.ac.in

Draft — October 21, 2026

Abstract

Objective. EEG classification of Alzheimer's disease (AD) and frontotemporal dementia (FTD) on one public dataset reports accuracies from about 50% to 99%, but studies change representation, graph, model and evaluation protocol at once. We isolated representation and graph construction under a fixed graph neural network, and tested a blood–brain-barrier-linked EEG marker. *Approach.* On OpenNeuro ds004504 (36 AD, 23 FTD, 29 controls) we crossed seven node-feature representations with three graph constructions under one compact, untuned graph convolutional network, on 25 frozen subject-level splits and 25 confirmation splits, with matched baselines, 16 ablations, leakage and permutation experiments, a simulation of the tests' calibration, a Parkinson's disease validation cohort ($n = 149$), and paroxysmal slow-wave events (PSWEs). *Main results.* Descriptively, representation outweighed graph construction in every analysis. In simulation, uncorrected and subject-level tests flagged differences between equally informative representations in 34–80% of datasets; with a calibrated adjustment for overlapping training sets the representation effect was significant externally ($p = 0.031$) but not on ds004504 ($p = 0.258$). Graph construction had a negligible effect (graph types equivalent within ± 0.05 macro-F1 on average), and a learned graph did no better. The best configuration (wavelets, spatial graph) reached macro-F1 0.596 (corrected 95% CI 0.487–0.706); non-graph and deep-learning baselines did not differ significantly. Window-level splitting raised macro-F1 to 0.983. PSWEs were more frequent in AD (rate ratio 2.96) and FTD (2.07) but not in Parkinson's disease; the excess tracked background slowing ($\rho = 0.88$) and was not detectable among patients whose slowing overlapped with controls'. *Significance.* The evaluation protocol moves reported performance far more than any modelling choice we tested; with 88 subjects, representation differences smaller than about 0.19 macro-F1 cannot be established once training-set overlap is accounted for. Fixed-threshold PSWEs on scalp EEG mainly index background slowing. Code, folds and raw runs are public.

Keywords: EEG; graph neural network; Alzheimer's disease; frontotemporal dementia; signal representation; functional connectivity; blood–brain barrier; paroxysmal slow-wave events; data leakage; cross-validation; statistical inference; reproducibility

1 Introduction

Dementia affected an estimated 57 million people worldwide in 2019, a number forecast to reach 153 million by 2050 [1]. Alzheimer's disease (AD) is its most common cause, accounting for an estimated 60–80% of cases [2], while frontotemporal dementia (FTD) weighs disproportionately on younger people: it accounts for about 3% of dementia diagnosed after 65 but about 10% before [3], an age

group in which some 3.9 million people live with dementia worldwide [4]. Telling AD and FTD apart is clinically consequential, not least because the first disease-modifying anti-amyloid therapies are indicated only for early AD with confirmed amyloid pathology [5]. Scalp electroencephalography (EEG) is attractive for this task because it is cheap, non-invasive and available in routine care, in contrast to amyloid PET or MRI. Its physiological basis is well established: AD is increasingly understood as a disorder of large-scale brain networks [6], and graph-theoretical EEG studies have shown that AD and FTD disturb those networks in *different* ways [7], which is what makes a three-class AD/FTD/control problem scientifically meaningful rather than merely harder.

Machine-learning work on this question has moved from spectral markers and handcrafted features with classical classifiers [8] to convolutional networks and transformers [9, 10] and, most recently, to graph neural networks (GNNs) that treat electrodes as nodes and inter-electrode coupling as edges [11–17]. The public OpenNeuro dataset ds004504 [18, 19], with 88 subjects recorded on one device under one protocol, has become the de-facto benchmark. Reported accuracies on this single dataset span roughly 50–99%, from 54% for the three-class problem under leave-subjects-out validation [20] to 99% under segment-level splits [21].

Part of that spread reflects different tasks (two or three classes), but much of it comes from how models are evaluated, and it hides a methodological problem. Published studies change several things at once: (i) the *signal representation* fed to the model, (ii) the *graph construction* that defines which electrodes are connected, (iii) the *model architecture*, and (iv) the *evaluation protocol*. Independent evidence suggests that the last of these matters most. A 2026 scoping review of this exact benchmark concluded that apparent deep-learning gains "arise not from inherent representational advantage, but from data leakage" [22], echoing a broader warning about leakage across machine-learning-based science [23] and in neuroimaging prediction specifically [24]. In clinical machine learning, splitting records rather than subjects has long been known to overestimate accuracy for new patients [25], and segment-level splitting is common in, and inflates, deep-learning studies of translational EEG [26, 27]. A 2026 comparison of five time–frequency representations under each of three pretrained convolutional networks [21] isolated representation, but its authors note that it "does not explicitly model the true spatial topology" and suggest that graph-based or topology-aware schemes may improve on it; others observe that most studies "overlook spatial connectivity between channels" [15], and explicitly call for graph convolutional networks in future work [28, 29]. No study has varied representation and graph construction together while holding the model and the evaluation protocol fixed.

A second, biologically motivated gap concerns the blood–brain barrier (BBB). BBB breakdown is an early biomarker of cognitive dysfunction [30, 31], and Milikovsky et al. described *paroxysmal slow-wave events* (PSWEs) — episodes in which the EEG median power frequency falls below 6 Hz for at least 5 s — in patients with AD, linked them to cognitive impairment, localised them to cortex with BBB dysfunction in epilepsy, and induced them in rodents by albumin extravasation [32, 33]. PSWEs have since been examined in Parkinson's disease [34], epilepsy [35], obstructive sleep apnoea [36] and healthy ageing [37], but, to our knowledge, never in FTD and never on an open dementia cohort. Whether this BBB-associated EEG marker carries diagnostic information beyond ordinary EEG slowing has not been tested.

We therefore ran a controlled factorial experiment. Seven node-feature representations (spectral, connectivity, wavelet, time-domain, raw, fused, and PSWE-derived) were crossed with three graph constructions (spatial, functional, hybrid) under a single fixed two-layer graph convolutional network [38], with subject-level repeated cross-validation, identical frozen folds for every condition, and hyperparameters fixed before any result was seen. Because the model is the control variable rather than the contribution, it is deliberately small. Around this core we measured the inflation produced by window-level (leaky) splitting on the identical pipeline, quantified age/sex confounding, and

characterised PSWE burden across AD, FTD and controls.

Our contributions are:

1. **Methodological.** A fully crossed 7×3 protocol that separates signal representation, graph topology and model architecture, with primary tests adjusted for overlapping training sets and calibrated by simulation, paired statistics over 25 frozen subject-level test splits, an independent set of confirmation splits that removes selection bias, and ablations that test whether the graph matters at all.

2. **Empirical.** Effect sizes, confidence intervals and equivalence tests for representation and graph construction in three-class AD/FTD/control classification, with non-graph twins and a demographic confound floor as reference points — and the finding that, with 88 subjects, neither factor's effect can be established once training-set overlap is accounted for.

3. **Statistical.** A simulation of the study's own resampling showing that tests which treat cross-validation splits or held-out subjects as independent flag differences between equally informative representations in 34–80% of datasets, and a Nadeau–Bengio-type adjustment that controls this in that simulation; other benchmark comparisons of this kind would need the same check before relying on it.

4. **Protocol.** A quantification, with the model and data held fixed, of how much window-level splitting inflates results on this benchmark and on a second cohort — not a new finding [25–27], but a calibration factor for reading published numbers on ds004504 — and an external test of the whole protocol on a Parkinson's disease cohort.

5. **Biological.** The first characterisation of BBB-associated PSWEs in FTD, an independent-cohort test of the AD finding, a replication in Parkinson's disease, and a test of whether the marker is separable from background slowing.

6. **Reproducibility.** Code, seeds, frozen folds and every raw run released; the full study re-runs on free cloud resources, on CPUs except for 0.57 GPU-hours for the deep-learning baselines.

2 Related work

Signal representations for EEG dementia classification. Spectral representations dominate: relative band power and EEG slowing are the most replicated AD markers [9, 39, 40]. Time-domain and complexity features (Hjorth parameters, entropies) are cheap and interpretable [10, 40]. Time–frequency transforms such as the continuous or discrete wavelet transform retain non-stationary structure that a power spectrum averages away [21, 41], and several recent studies encode EEG as images for convolutional networks [16, 21]. Representation choice is thus treated as a first-class question — but almost always with convolutional models, and without varying the spatial structure the model sees.

Graph neural networks for EEG. Spectral graph convolution [38] and attention-based variants [42] underpin most EEG GNNs, reviewed in [11]. For dementia, GNNs have been built on functional-connectivity graphs [12], learned adjacency [13], multi-connectivity inputs [17], time-resolved graphs [16] and spatial–functional hybrids [15]; for Parkinson's disease, interpretable graph learning has been used to localise abnormal connectivity [29]. The closest precedent [12] compared connectivity measures as GNN inputs but held node features fixed. A 2025 preprint applied a graph transformer

to ds004504 for FTD versus controls with window-level prediction [43]. None of these studies crosses node-feature representation with edge definition under a fixed model.

Graph construction for EEG. Edges may be defined spatially (inter-electrode distance), functionally (statistical coupling), by learning, or by combinations of these. Phase-lag-based measures [44, 45] suppress the zero-lag coupling produced by volume conduction, which otherwise inflates functional graphs; weighted graphs are recommended over binarised ones for EEG [7], and phase-synchronisation networks differ between dementia subtypes [41, 46]. Spatial graphs carry no subject-specific information on their own and act as a smoothing prior; whether that prior helps is an empirical question we test directly.

Evaluation protocol and leakage. Leakage — information from test data reaching training — is a general failure mode of machine-learning-based science [23] and inflates prediction in connectome-based neuroimaging, more so in small datasets [24]. When the use case is prediction for new individuals, record-wise cross-validation greatly overestimates accuracy [25]. In EEG, windows cut from one recording are highly correlated, so splitting *windows* rather than *subjects* lets a model recognise individuals instead of disease; improper partitioning can understate cross-participant error rates severalfold [47], segment-based holdout overestimated performance in an Alzheimer's dataset and was used by most translational deep-learning EEG studies surveyed [26], and an audit of ADHD and Parkinson's cohorts — the latter our external cohort — reproduced the inflation [27]. The 2026 scoping review of ds004504 attributes the upper end of reported accuracies to the same effect [22]; one study on an 88-subject AD/FTD/control cohort reported 98.8% segment-level accuracy with random splitting and 84.3% with subject-wise splitting [21]. Best-practice guidance [48] and the large error bars of cross-validation in small samples [49] add a separate caution for cohorts of our size.

EEG and the blood–brain barrier. BBB breakdown precedes and predicts cognitive decline [30, 31]. Milikovsky et al. proposed PSWEs as an EEG manifestation of BBB-related cortical dysfunction in AD and epilepsy [32], with a companion study tracing the mechanism to albumin-activated TGFβ signalling [33]. Subsequent work found PSWEs uncommon in Parkinson's disease [34] and explored them as an epilepsy biomarker [35]; they have also been linked to cognitive impairment in obstructive sleep apnoea [36] and studied with MEG in healthy adults in relation to age and cognitive performance [37]. Because ds004504 includes no BBB imaging, we treat PSWEs strictly as a *BBB-dysfunction-associated* EEG marker, not as a measurement of BBB integrity.

Positioning. Table 1 summarises how this study differs from the closest prior work.

Table 1: Positioning against the closest prior work on EEG dementia classification, verified against each paper's full text. ✓ = yes; (✓) = partly (see notes); – = no.

Study	Varies repr.	Varies graph	Fixed model	Subject-level split	Public data
DICE-Net [9]	–	–	✓	✓	✓
FC-measure GNN [12]	–	✓	(✓)[a]	✓	–[e]
Adaptive GCN [13]	–	(✓)[b]	–	✓	–[e]
TFR comparison [21]	✓	–	(✓)[c]	(✓)[d]	✓
Spatial–functional GNN [15]	–	–	–	✓	✓
Multi-connectivity GCN [17]	–	(✓)[b]	–	✓	✓
Graph transformer [43]	–	–	✓	✓	✓
This study	✓	✓	✓	✓	✓

[a] Same GNN architecture for every connectivity measure, but hyperparameters tuned separately for each. [b] Graph construction varied only by ablating the proposed model's graph module (learned graph replaced by a fully connected one [13]; one of two connectivity matrices removed [17]). [c] Representations compared under each of three pretrained CNNs (MobileNetV2, ResNet-50, InceptionV3). [d] Reports both random (segment-level) and subject-wise splits (98.8% versus 84.3% accuracy). [e] In-house clinical cohort (20 AD, 20 controls); no public release stated.

3 Materials and methods

Figure 1 summarises the design.

3.1 Dataset

We used OpenNeuro ds004504 [18, 19]: resting-state, eyes-closed scalp EEG from 88 participants recorded at AHEPA University Hospital (36 AD, 23 FTD, 29 cognitively normal controls, CN). All values below were read from the dataset's own BIDS files rather than from the descriptor paper. Recordings use 19 electrodes of the 10–20 system (Fp1, Fp2, F3, F4, C3, C4, P3, P4, O1, O2, F7, F8, T3, T4, T5, T6, Fz, Cz, Pz) referenced to A1–A2 according to the BIDS sidecar (the dataset README describes a Cz reference; we re-reference to the average, so this does not affect any analysis), sampled at 500 Hz on a Nihon Kohden EEG 2100 system, with a 0.4–50 Hz acquisition filter. Recording length per subject ranges from 307 to 1291 s (mean 802 s; 19.6 h in total). Group demographics are given in Table 2. The data are released under CC0; the study uses only de-identified public data. The dataset's own metadata cite a version-1.0.9 DOI that did not resolve when checked; we cite version 1.0.8, whose EEG data are identical (the 1.0.9 changelog records only a README edit).

3.2 Fixed preprocessing

Every condition shares one code path, implemented with MNE-Python [50]. Raw recordings were band-pass filtered at 0.5–45 Hz, notch filtered at 50 Hz, and re-referenced to the common average. Independent component analysis (extended Infomax [51], 18 components — one fewer than the channel count because of the average reference) was fitted on a 1 Hz high-passed copy of each recording, and components that ICLabel [52] assigned to eye or muscle artefact with probability ≥ 0.8 were removed. Data were then resampled to 128 Hz and cut into non-overlapping 10 s windows (1,280 samples), giving 7,013 windows in total (30–129 per subject). ICA is fitted per recording, without labels or other subjects' data, and therefore cannot leak information across cross-validation folds; every transform that is estimated across subjects (feature scaling, demographic residualisation) is fitted on training subjects only.

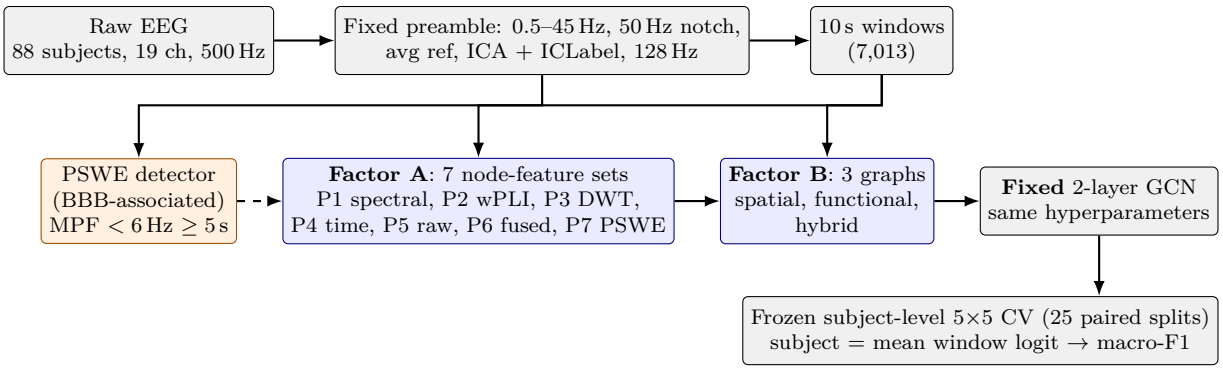


Figure 1: Study design. Grey components are held fixed; only the node features (Factor A) and the graph construction (Factor B) vary, in a fully crossed 7×3 design evaluated on the same 25 frozen subject-level splits. The PSWE detector (orange) supplies the blood–brain-barrier-associated marker, used as P7 node features and, in an ablation, as a co-occurrence graph.

Table 2: Participant characteristics (mean ± SD), from `participants.tsv`. MMSE is reported for description only and is never a model input. The dataset README gives the FTD MMSE as 22.17 ± 8.22; the SD computed from the participant file is the one shown.

	AD ($n = 36$)	FTD ($n = 23$)	CN ($n = 29$)
Age (years)	66.4 ± 7.9	63.7 ± 8.2	67.9 ± 5.4
Sex (F/M)	24/12	9/14	11/18
MMSE	17.8 ± 4.5	22.2 ± 2.6	30.0 ± 0.0

3.3 Factor A: node-feature representations

Each window is a graph whose 19 nodes are the electrodes. Seven representations define the node features $X \in \mathbb{R}^{19 \times F}$:

P1, spectral ($F = 7$). Relative power in five bands (delta 0.5–4, theta 4–8, alpha 8–13, beta 13–30, gamma 30–45 Hz) from Welch's method [53] (2 s Hann segments, 50% overlap), normalised spectral entropy, and peak frequency within 4–15 Hz.

P2, connectivity ($F = 5$). Node strength of the weighted phase-lag index (wPLI) [44, 45] in each band: the mean wPLI between a channel and the other 18.

P3, wavelet ($F = 20$). Five-level Daubechies-4 discrete wavelet transform [54]; for the sub-bands A5, D5, D4, D3 and D2 (D1, 32–64 Hz, lies above the low-pass cut-off), relative energy, normalised coefficient entropy, mean absolute coefficient and standard deviation.

P4, time domain ($F = 10$). Mean, variance, root mean square, skewness, kurtosis, zero-crossing rate, Hjorth mobility and complexity [55], sample entropy [56] ($m = 2$, $r = 0.2\,\sigma$) and line length. Line length replaces Hjorth activity, which equals the variance.

P5, raw ($F = 1280$). The window z-scored per channel.

P6, fused ($F = 34$). P1, P3 and P2 concatenated, plus each node's weighted clustering coefficient [57] and eigenvector centrality in the alpha-band wPLI graph.

P7, PSWE ($F = 4$). From the PSWE detector (Section 3.8): fraction of the window spent in a PSWE, mean and minimum median power frequency in the window, and the local PSWE onset rate within ±60 s.

The wPLI between channels i and j was computed within each window from band-passed (4th-order Butterworth, zero-phase) analytic signals $z_i(t)$ as

$$\mathrm{wPLI}_{ij} = \frac{\left|\langle \mathrm{Im}(z_i z_j^*)\rangle_t\right|}{\langle \left|\mathrm{Im}(z_i z_j^*)\right|\rangle_t}, \tag{1}$$

which is insensitive to zero-lag coupling from volume conduction [45].

3.4 Factor B: graph construction

Three edge definitions were crossed with every representation. *Spatial*: a Gaussian kernel $w_{ij} = \exp(-d_{ij}^2/2\sigma^2)$ on Euclidean distances between standard 10–20 electrode positions, with σ the median pairwise distance; this graph is identical for every window and subject. *Functional*: the alpha-band wPLI of the window. *Hybrid*: the element-wise product of the two. Each graph was sparsified by keeping the $k = 4$ strongest edges per node, symmetrised as $A = \max(A, A^\top)$, kept weighted, and normalised as $\hat{A} = \tilde{D}^{-1/2}(A + I)\tilde{D}^{-1/2}$.

3.5 The fixed model

The model is the control variable, not the contribution, so a single small graph convolutional network (GCN) [38] was used everywhere: $H_1 = \mathrm{Dropout}_{0.5}(\mathrm{BN}(\mathrm{ReLU}(\hat{A}XW_0)))$, $H_2 = \mathrm{ReLU}(\hat{A}H_1W_1)$, mean pooling over the 19 nodes, and a linear layer to three classes, with $W_0 \in \mathbb{R}^{F\times 64}$ and $W_1 \in \mathbb{R}^{64\times 32}$. This gives 2,819 parameters for $F = 7$ and 3,651 for $F = 20$; P5 is the unavoidable exception (84,355 parameters, because $F = 1280$). Two hops of propagation already reach most of a sparsified 19-node graph, and deeper GCNs risk over-smoothing; depth is tested as an ablation. Graph attention networks [42] learn edge weights and would partly undo the edge definitions being compared, so a GAT is used only as a sensitivity check.

Training used AdamW (learning rate 10^{-3}, weight decay 10^{-4}), batch size 64, and at most 200 epochs with early stopping (patience 20) on validation macro-F1 computed at subject level. The loss was weighted by inverse class frequency (over subjects) divided by each subject's window count, so every subject contributes equally regardless of recording length. A subject's prediction is the mean of its window logits. Features were standardised with statistics from training windows only. All hyperparameters were fixed before any result was seen and are identical in every condition.

3.6 Evaluation protocol

We used subject-level, stratified, repeated 5-fold cross-validation with five repeats (seeds 0–4), for 25 test splits of 17–18 subjects each. Within each training set, a stratified 20% of subjects formed the early-stopping validation set; it was never used to tune hyperparameters. The fold assignment was generated once, verified by assertions that training, validation and test subjects are pairwise disjoint and that every subject is tested exactly once per repeat, and committed to the code repository before any training run. All conditions use the same 25 splits, so every comparison is paired. The full factorial comprises 7×3 cells $\times$ 25 splits = 525 training runs.

Confirmation folds. Choosing the best cell, the cells to ablate and the cell to permute on the same splits that then estimate their performance biases those estimates upwards (the winner's curse). We therefore generated a second, independent fold assignment with the same procedure (seeds 5–9; again 25 splits) and committed it to the repository after the main grid had been analysed but before any run used it. The whole grid was re-run on these folds (525 runs). The cell ranked first on seeds 0–4 is reported with its score on seeds 5–9 as a selection-free estimate, and all ablations, the leakage experiment and the permutation test reported in the main text use the confirmation folds. Their selection-fold counterparts are given in the Supplementary material.

The primary metric is subject-level macro-averaged F1. We also report balanced accuracy, one-vs-rest macro ROC-AUC, PR-AUC, per-class sensitivity and specificity, Cohen's κ, the Brier score, expected calibration error and confusion matrices. Plain accuracy is reported but never used for ranking.

3.7 Baselines, ablations and controls

Baselines. A logistic regression on age and sex alone (the demographic confound floor); a stratified random classifier (chance); and logistic regression, RBF-kernel SVM and random forest trained on the flattened node features of each representation with the same folds, weighting and subject aggregation as the GCN (a non-graph twin of every representation). The MMSE is never used as an input: control participants all score 30, so it is a proxy for the label.

Deep-learning baselines. EEGNet-8,2 [58] and a compact 1D-CNN were trained on the raw windows with the same splits, training rule and aggregation as the GCN (Supplementary methods).

Ablations. Run on the two best cells of the main grid, chosen automatically from the seed 0–4 results and evaluated on the 25 confirmation splits, each paired with the same cell's confirmation-grid run: no graph (identity adjacency, i.e. a per-node MLP); degree-preserving random graphs; $k \in \{3, 6\}$ and a fully connected graph; binarised edges; depth 1 and 3; GAT instead of GCN; no subject-equal weighting; demographic residualisation (age and sex regressed out of every feature, coefficients fitted on training windows); adding P7 (PSWE) features to the best representation; the PSWE co-occurrence graph (Jaccard overlap of channels' PSWE time) as the edge type; 5 s and 20 s instead of 10 s windows (A7; the preprocessed recordings re-windowed and every feature and graph rebuilt); and preprocessing without ICA/ICLabel artefact removal (A8; every other step unchanged). Holm correction is applied within each cell's family of 16 ablations.

Leakage. The best cells were re-run with *window-level* stratified splitting — windows assigned to folds at random, ignoring subject identity (record-wise cross-validation in the terminology of [25]; segment-based holdout in [26]) — with everything else unchanged.

Permutation test. For the best cell, subject labels were permuted (all windows of a subject keep one label) and the first repeat's five folds re-run for each permutation; the p-value is $(1 + \#\{\text{null} \geq \text{observed}\})/(1 + N)$. The primary test uses 100 permutations on the first confirmation repeat (seed 5); the original 200 permutations on the first selection repeat (seed 0) are reported for completeness.

Stopping-rule sensitivity. The grid was re-run with patience 50 and without early stopping (Supplementary methods).

Severity (added). Results stratified by MMSE (analysis only) are described in the Supplementary methods.

Learned adjacency (added). To test whether a graph fitted to the data would help where the three fixed constructions did not, the GCN was given a trainable, symmetric, non-negative 19×19 adjacency (softplus of free parameters, symmetrised and degree-normalised), initialised at the spatial graph and learned jointly with the other weights (361 extra parameters). It was run for every representation on the 25 selection splits and for P3 on the 25 confirmation splits, and compared with the fixed spatial graph split by split; everything else was unchanged.

Occlusion importance (added). For the best cell (P3 $\times$ spatial), each of the 25 selection-split models was re-fitted with its original seed (the runs store predictions, not weights) and its test subjects were scored again with one electrode's features replaced by their training mean (zero after standardisation), or with one wavelet sub-band's four features replaced in the same way on every electrode; the model was not retrained for each occlusion. Importance is the drop in the probability assigned to each test subject's true class, averaged within subject over the five repeats and then over the 88 subjects, with a subject-bootstrap 95% confidence interval, and also within each class. Occlusion measures what a trained model relies on, not where the disease is. Re-fitting on the original feature cache reproduced the stored predictions of all 25 splits exactly (largest probability difference 0), so the occluded models are the models of the main grid.

3.8 PSWE detection

Following Milikovsky et al. [32], with parameters as reported in [34], we computed the median power frequency (MPF) of each channel from FFTs of 2 s Hann-windowed segments with 1 s overlap (one MPF value per second) over 1–45 Hz. The frequency range is not reported in the source and was set to our pass-band. A PSWE is a run of at least five consecutive seconds with MPF < 6 Hz. Detection is per channel and per recording, and uses no labels. For each subject we computed the PSWE rate (events per minute, averaged over channels) and the fraction of time spent in PSWEs.

Background-relative PSWE (exploratory, post hoc). Events defined as drops below each channel's own median MPF are described in the Supplementary methods.

Detector robustness and collinearity. We repeated the PSWE statistics with zero-padded spectra, artefact rejection, without ICA and over a grid of event rules (A11), and examined the collinearity of the slowing adjustment (Supplementary methods).

3.9 External validation cohort

To test whether the conclusions depend on one cohort, the identical protocol was applied to OpenNeuro ds004584 [59]: 149 participants (100 with Parkinson's disease, 49 controls), 2 min to 5 min of eyes-open resting EEG, 63 channels at 500 Hz, with patients on their usual levodopa [60]. Only what must differ between datasets was changed, in a configuration overlay: the 19 electrodes of the main montage were selected in the same node order, the online reference (Pz), which is not stored, was restored as a flat channel before average re-referencing (exact, because every stored channel is referenced to it); the mains notch was 60 Hz; the task is two-class; and the Montreal Cognitive Assessment replaced the MMSE as the analysis-only cognitive score (the UPDRS, a Parkinson's-only label proxy, is never used). Features, graphs, model, training, cross-validation

(seeds 0–4, frozen before any training on this cohort) and statistics were unchanged. Questions were fixed before the first run: whether the ranking of the 21 cells and of the 7 representations transfers (Spearman and Kendall correlations of cell means); how the configuration selected on ds004504 performs without re-selection (with 100 label permutations); whether matched-training non-graph twins, the age-and-sex floor and window-level leakage behave as on ds004504; and whether fixed-threshold PSWEs are more frequent in Parkinson's disease and track slowing, as a test of the earlier report that they are uncommon in this disease [34]. The two datasets differ in site, device, disease, eyes-open versus eyes-closed recording and recording length, so absolute scores are not compared.

3.10 Statistical analysis

Primary inference. The planned unit of analysis is the test split. The primary test of each factor (representation, graph construction and their interaction) is the two-way repeated-measures ANOVA on macro-F1 over the 25 splits, with the Nadeau–Bengio variance inflation applied to each contrast to allow for the overlap between training sets, $F_{\mathrm{NB}} = F/(1 + Jn_{\mathrm{test}}/n_{\mathrm{train}})$ with $J = 25$ (a factor of 8.8), on the same degrees of freedom with Greenhouse–Geisser correction. For two levels this is exactly the squared Nadeau–Bengio t statistic [61], which is also used for every pairwise comparison (Holm-corrected). Tests that treat splits, or held-out subjects, as independent describe this cohort rather than the population; the adjustment bounds that optimism, and its calibration was checked by simulation.

By-subject mixed model (descriptive). To describe how individual subjects respond to each factor, we also treat the held-out subject as the unit. For every run we take the probability the model assigns to each test subject's true class and average it over the five repeats (each from a different partition), giving one value per subject and cell (1,848 values from 9,240 subject-level predictions for ds004504). These were modelled with a linear mixed model with representation, graph construction and their interaction as fixed effects (sum-to-zero coding), true class as a covariate, and random intercepts for subject, subject × representation and subject × graph, fitted by REML in statsmodels. The subject-by-factor terms matter: a subject's advantage under one representation recurs in every repeat and graph type, so omitting them treats thousands of correlated predictions as independent (a first version without them gave implausibly small p-values and was discarded). Each term is tested with a Wald F test against its own subject-level stratum (containment degrees of freedom); for balanced data the model reduces to a by-subject repeated-measures ANOVA, which our implementation reproduces exactly in tests. Its false-positive rate is nominal when the representations do not differ within this cohort (simulated subject-level data with subject-by-representation variation), but it treats the trained models as fixed: a representation that happens to suit this particular sample counts as a real difference. The calibration simulation below shows that, for claims about new samples, it then declares differences far too often, so we report it descriptively. We report model-based marginal means with 95% confidence intervals and Holm-corrected pairwise contrasts between representations. The same model was fitted to the confirmation folds and to the external cohort. Two cross-checks were added after review: a by-subject repeated-measures ANOVA on the same subject-by-cell values (Greenhouse–Geisser corrected), and the same mixed model fitted to the proportion of repeats in which each subject was *correctly* classified — a linear-probability approximation to a logistic mixed model of the binary decisions, which tests the decisions a clinician would see rather than the model's confidence.

Calibration of the tests (added). Because the adjusted F is our generalisation, we checked every test by simulation under the study's own resampling: 88 simulated subjects in the cohort's class sizes, 5×5 subject-level cross-validation with 20% of each training set held out, and seven 'representations' that are noisy views of one latent class signal, classified by logistic regression. Under the null all seven are equally informative in the population; under the alternatives their noise levels differ. For each of 200–300 simulated datasets per setting we applied the split-level repeated-measures F (uncorrected and adjusted) and the by-subject repeated-measures F on averaged true-class probabilities, to which the by-subject mixed model reduces for a balanced design. Because real representations are transforms of the same EEG, the null and the strongest alternative were repeated with the seven representations' noise correlated at 0.5 and 0.8.

Out-of-bootstrap comparison (added). As a check that does not rely on the adjustment, we resampled the cohort itself: in each of 200 replicates, subjects were drawn with replacement within each class, the GCN (spatial graph) was trained on the distinct drawn subjects (20% held out for early stopping) and scored on the subjects never drawn, for all seven representations on the same replicate. Each paired difference between representations was tested with its bootstrap standard error (Holm-corrected over the 21 pairs), and its 95% percentile interval is reported.

Descriptive analyses on macro-F1. For comparability with the literature we also report the unadjusted two-way repeated-measures ANOVA on macro-F1 (partial η^2, Greenhouse–Geisser-corrected p-values), which shares the mixed model's limitation. Because cross-validation scores share training data and are therefore not independent, every pairwise comparison is reported with both the Wilcoxon signed-rank test and the Nadeau–Bengio corrected resampled t-test [61], with Holm correction [62] within each family, Cohen's d and the matched-pairs rank-biserial correlation. All 21 cells were also compared with a Friedman test and a Nemenyi critical-difference analysis [63]. Cell-level 95% confidence intervals come from a subject-level bootstrap (2,000 resamples of subjects within each repeat) and from a t-interval over the 25 splits. The Nadeau–Bengio correction uses the ratio of test to training subjects of the models being compared: 17.6/56.3 for models fitted on the training subjects only (the main grid and the ablations) and 17.6/70.4 when validation subjects are added (the train + val regime).

Equivalence. A non-significant difference is not evidence of no difference. For every paired comparison read as "no advantage" (GCN versus its non-graph twin, graph types, ablations, deep-learning baselines, stopping rules and the learned adjacency) we report two one-sided tests (TOST) against a margin of ± 0.05 macro-F1 — about one additional correctly classified subject in a test fold of 17–18 — using the Nadeau–Bengio corrected standard error, with Holm correction within each family, together with the 90% confidence interval and the minimum difference detectable with 80% power. With the corrected standard error these tests are conservative: a comparison is called equivalent only if its whole corrected 90% interval lies within the margin; the calibration simulation confirms that tests built on this standard error do not over-reject, so equivalence and non-difference conclusions are, if anything, conservative.

For PSWEs, group differences were tested with Kruskal–Wallis and Dunn post-hoc tests (Holm-corrected, with rank-biserial effect sizes), and with negative-binomial regression of event counts using recording length as exposure, adjusted for age and sex and, in further models, for relative delta+theta power and background MPF. The relation to slowing was assessed with Spearman and partial (age- and sex-adjusted) rank correlations, and the relation to cognition with Spearman correlation against MMSE within patients only. All analyses were run with SciPy, statsmodels and custom code that is tested against reference implementations.

3.11 Deviations from the pre-specified plan

The study plan (factorial design, preprocessing, model, training, folds, baselines, ablations A1–A12, permutation test, leakage experiment, deep-learning baselines, external validation and the fixed PSWE definition) was committed to the code repository before any data were analysed; all changes are logged there with dates. This time-stamped commit is a pre-specification, not a registration with an external registry, so it relies on the public commit history rather than independent oversight; Table S1 gives the status of every analysis. The following were decided after results were seen. (i) *Primary analysis.* The plan named the split-level repeated-measures ANOVA on macro-F1 as primary; it remains primary, with the Nadeau–Bengio adjustment for overlapping training sets added after review. An internal review first made a by-subject mixed model primary (a version without subject-by-factor random effects gave implausibly small p-values and was discarded); when the calibration simulation showed that this model, like the unadjusted ANOVA, does not control false positives for claims about new samples, it was reported descriptively instead (Section 3.10). (ii) *Selection bias.* The confirmation folds (seeds 5–9) replaced the planned ten-repeat re-run of the top three cells; they were generated after the main grid was analysed and committed before any run used them, and ablations, leakage and permutation tests were moved onto them. (iii) *Matched training data.* Classical baselines were first fitted on training plus validation subjects; we added train-only fits and train + val GCN refits, and corrected the Nadeau–Bengio ratio. (iv) *Added analyses.* The background-relative PSWE definitions (Section 3.8), the stopping-rule sensitivity, the PSWE detector robustness and collinearity checks, the equivalence tests, the learned-adjacency GCN, occlusion importance, the two mixed-model cross-checks, the severity-stratified analysis, an EEGNet run with dropout 0.25, the Nadeau–Bengio-adjusted omnibus tests, the graph-type equivalence tests within representations, the calibration simulations (including one matched to the data), the re-fit checks and the out-of-bootstrap comparison were added after the corresponding results were known; their definitions, variants and margins were fixed in the configuration before they were run. A7, A8 and the A11 PSWE rule grid were planned but run late. (v) *Changed or not done.* Planned integrated-gradients explanations were replaced by occlusion, which is model-agnostic and reads directly as lost probability for the true class; the planned connectome and embedding plots and calibration diagrams were not carried out. (vi) *Implementation details.* Hjorth activity in P4 was replaced by line length (it duplicates the variance); PSWEs are detected per channel rather than over averaged regions, with an MPF range set by us (the source does not report one); EEGNet's max-norm constraints were replaced by the shared weight decay and its dropout kept at the shared 0.5 (a run with the cross-subject value of 0.25 is reported as a sensitivity analysis).

3.12 Compute

All computation ran on free Kaggle sessions; resources are detailed in the Supplementary methods.

4 Results

4.1 Data quality and the confound floor

The fixed preprocessing produced 7,013 windows from all 88 subjects (Figure S9); ICA removed on average 2.4 components per recording. The expected EEG slowing in dementia was present before any model was trained: the median theta/alpha power ratio was 1.57 in AD, 0.68 in FTD and 0.28 in controls, with the posterior alpha of controls replaced by diffuse delta and theta in AD (Figure S10).

A logistic regression on age and sex alone reached a macro-F1 of 0.411, against 0.302 for a stratified random classifier. Demographics therefore carry some group information, consistent with

the sex imbalance between groups (Table 2). We use 0.411, not 1/3, as the floor that EEG-based models must clear.

4.2 Representation versus graph construction

Across the 21 cells on the selection folds (525 runs; Table S8), the best configuration was P3 × spatial, with subject-level macro-F1 0.586 ± 0.112 (subject-bootstrap 95% CI 0.538–0.628). Because this cell was chosen as the maximum of 21, that figure is optimistically biased. Re-evaluated on the independent confirmation folds, the same cell scored 0.596 ± 0.089 (95% CI 0.487–0.706 with the Nadeau–Bengio correction for overlapping training sets; subject-bootstrap interval, which describes this cohort only, 0.554–0.643; balanced accuracy 0.610, macro ROC-AUC 0.741, Cohen's κ 0.428) and again ranked first of 21. The ordering of all cells was largely preserved (Spearman $\rho = 0.86$ between the two fold sets; Figure S3). We therefore take 0.596 as the selection-free estimate for the best configuration.

Table 3 lists every test of the factorial, the primary one first. In the primary analysis — the split-level test adjusted for overlapping training sets — neither representation ($F_{\mathrm{NB}} = 1.34$, $p = 0.258$; confirmation 1.54, $p = 0.188$) nor graph construction ($p = 0.629$) nor their interaction ($p = 0.993$) was significant on ds004504; in the external cohort representation was ($p = 0.031$). The unadjusted split-level ANOVA and the by-subject mixed model, which treat splits or subjects as independent, found strong representation effects in all three analyses, and the mixed model also a graph effect that spans only 0.007 in the probability assigned to the true class; their details are descriptive and given in the Supplementary results (Table S2, Figure S1). The last row of Table 3 gives each test's false-positive rate in the calibration simulation, examined in Section 4.3.

On macro-F1 averaged over representations, every pairwise comparison of graph types was formally equivalent within ±0.05 on both fold sets (TOST with the Nadeau–Bengio correction, largest Holm-corrected $p = 0.027$ after correction; all corrected 90% intervals within ±0.043; Table S12). This is positive evidence that, on average, the choice of graph does not matter at that resolution, not merely an absence of significance. Within single representations the same comparisons were inconclusive: none of the 42 was shown equivalent or different (median detectable difference 0.12).

On macro-F1 with the 25 splits as units and without the adjustment (descriptive), representation had a large effect on both fold sets (Table S9; selection: $F(6, 144) = 11.83$, Greenhouse–Geisser $p < 0.001$, partial $\eta^2 = 0.33$; confirmation: $F = 13.57$, $p < 0.001$, $\eta^2 = 0.36$). Averaged over graph types, the wavelet (P3, 0.569; confirmation 0.583), spectral (P1, 0.552; 0.528) and fused (P6, 0.543; 0.541) representations formed a leading group; raw windows (P5, 0.489) were intermediate; and connectivity strength (P2, 0.453), PSWE features (P7, 0.448) and time-domain statistics (P4, 0.434) trailed (Figure S11). Graph construction had a small effect on the selection folds ($F(2, 48) = 4.03$, $p = 0.026$, partial $\eta^2 = 0.14$; marginal macro-F1 0.506 spatial, 0.497 functional, 0.492 hybrid; Figure S11) that did not replicate on the confirmation folds ($F = 0.03$, $p = 0.957$; marginal means 0.490, 0.491 and 0.490). The representation × graph interaction was not significant on either (selection $F(12, 288) = 1.43$, $p = 0.194$; confirmation $F = 1.58$, $p = 0.149$; Figure 2). Descriptively, a Friedman test across all cells was significant ($p < 0.001$), but the Nemenyi critical difference (6.26 ranks) left the top cells statistically indistinguishable from one another (Figure S12).

Pairwise comparisons depend on how the dependence between cross-validation splits is handled. With Holm-corrected Wilcoxon tests (descriptive; like the Friedman test they ignore that dependence), 9 of the 21 representation pairs differed on the selection folds and 10 on the confirmation folds; with the more conservative Nadeau–Bengio corrected t-test, 0 and 0 did. For example, P3 exceeded P4 by 0.135 macro-F1 on the selection folds (Cohen's $d = 1.50$; Wilcoxon–Holm $p < 0.001$; Nadeau–

Table 3: Every test of the factorial. Primary: the split-level repeated-measures ANOVA on macro-F1 with the Nadeau–Bengio adjustment for overlapping training sets (F_{NB}, Greenhouse–Geisser p). Descriptive: the same ANOVA unadjusted, and the by-subject mixed model of the true-class probability. Last row: rate at which each test declared a representation effect when seven representations were equally informative, in the calibration simulation.

		Primary (adjusted)		Descriptive p	
Data	Term	F_{NB} (df)	p	Split level	By subject
ds004504, selection	Representation	1.34 (6, 144)	0.258	<0.001	<0.001
	Graph	0.46 (2, 48)	0.629	0.026	<0.001
	Interaction	0.16 (12, 288)	0.993	0.194	<0.001
ds004504, confirmation	Representation	1.54 (6, 144)	0.188	<0.001	<0.001
	Graph	0.00 (2, 48)	0.994	0.957	0.002
	Interaction	0.18 (12, 288)	0.987	0.149	<0.001
ds004584	Representation	2.63 (6, 144)	0.031	<0.001	<0.001
	Graph	0.19 (2, 48)	0.824	0.207	0.002
	Interaction	0.23 (12, 288)	0.949	0.080	<0.001
False positives under the null (simulation)			0%	53–80%	34–43%

Bengio–Holm $p = 0.390$). The same caveat applies to the unadjusted omnibus tests, which also treat the splits as independent; with the adjustment (the primary analysis) no factor was significant.

4.3 Which tests can be trusted

The calibration simulation shows which of these tests can be trusted. When seven representations were equally informative in the population, the uncorrected split-level test declared a difference in 80% of simulated datasets and the by-subject test in 40%, because a representation that happens to suit this sample keeps its advantage under every re-split of the same subjects; the adjusted test did so in 0%. The spread of representation means under that null, 0.13 macro-F1, is about the spread we observed (0.13). The price of the adjustment is power: for true spreads of 0.22–0.23 it detected a difference in only 20–34% of datasets. Real representations are transforms of the same EEG, so their errors are correlated; with the simulated representations' noise correlated at 0.5 or 0.8, the null spread shrank to 0.08–0.11, below the observed 0.13, but the unadjusted tests still declared a difference in 53–74% (split level) and 34–43% (by subject) of null datasets, the adjusted test in 0%, and the adjusted test's power rose to 37–49%. These settings were illustrative; to judge our own spread we matched the null to the data, choosing the simulation's signal strength, noise and noise correlation so that it reproduced the observed mean macro-F1 (0.50 against 0.50), split-to-split SD (0.11 against 0.09) and between-representation correlation of split scores (0.16 against 0.19). In 300 matched null datasets the spread of representation means averaged 0.12 (95th percentile 0.19); the observed 0.13 lies well inside that distribution (empirical $p = 0.31$), and at the matched setting the unadjusted test still flagged 78% of null datasets against 0% for the adjusted one. The out-of-bootstrap comparison, which retrains on 200 resampled cohorts and does not depend on the adjustment, gave the same ordering (P3 highest at 0.548, P4 lowest at 0.422) but wide intervals: the largest difference, P3 minus P4, was +0.126 (95% percentile interval −0.073 to +0.287; Holm $p = 1.000$), the median bootstrap standard error of a paired difference was 0.095, and 0 of the 21 pairs differed. Each replicate trains on about 44 distinct subjects (against 56 in the cross-validation) and scores about 32, with the spatial graph only, so its scores are lower and its power smaller than

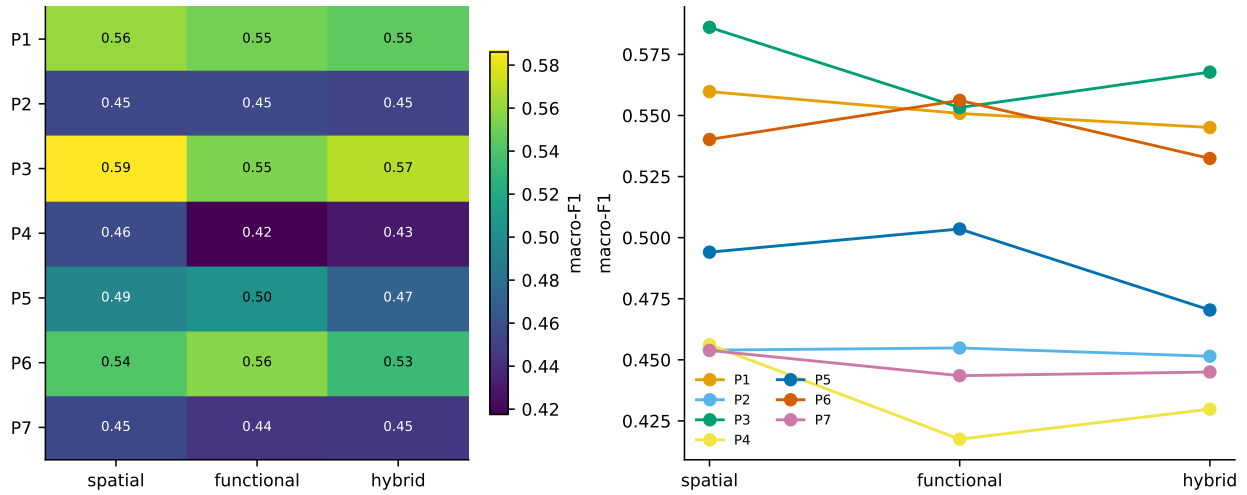


Figure 2: Mean macro-F1 of every representation × graph cell (left) and the same values as an interaction plot (right). Lines are close to parallel: P3 is best under the spatial and hybrid graphs and, under the functional graph, within 0.01 of P6; the leading group (P1, P3, P6) is the same for every graph type.

the cross-validation's; it is a check on the direction of the conclusion, not a more sensitive test. The significant subject-level and uncorrected tests above, and the replication of the ordering across fold sets of the same subjects ($\rho = 0.86$), therefore describe this cohort and do not by themselves establish a population difference; on ds004504 the representation effect is not established, and the calibrated evidence that representation matters comes from the external cohort, where the adjusted test was significant (Section 4.10).

4.4 What the models get right and wrong

For P3 × spatial, sensitivity was 0.771 for controls, 0.642 for AD and only 0.368 for FTD (specificities 0.786, 0.775 and 0.852). Across representations, controls were recognised best and FTD — the smallest group — worst; FTD recordings were most often misclassified as AD, except with PSWE features, where they were more often taken for controls (Figure S13). The ROC and precision–recall curves show the same ordering (Figure S14). The models were poorly calibrated (expected calibration error 0.265), so their probabilities should not be read as clinical risks.

Severity. Mild patients (MMSE ≥ 21) were not detectably harder to classify than moderate-to-severe patients (none of 8 tests significant; Supplementary results and Table S7).

What the best model relies on. Occluding one wavelet sub-band on every electrode cost far more than occluding any single electrode (Figure S8). The theta sub-band D4 (4–8 Hz) mattered most (drop in true-class probability 0.074, 95% CI 0.026–0.121), followed by D3 (8–16 Hz) (0.067); of the five bands, 3 had confidence intervals above zero, and the slowest, A5 (0–2 Hz), mattered least (0.019). The classes leaned on different bands: controls on D4 (4–8 Hz) (0.194), FTD on D2 (16–32 Hz) (0.177) and AD on D5 (2–4 Hz) (0.044). No single electrode was critical (largest drop 0.019; the four most important were Fp1, Fp2, F3, F4, and 12 of 19 had intervals above zero), so the information is spread across the scalp. That the frontopolar electrodes rank first may partly

reflect residual ocular activity after ICA, and these importances describe this model rather than the topography of disease.

4.5 Graph models do not beat their non-graph twins

Table S3 compares every representation's best GCN cell with classical classifiers trained on the same features, test splits and weighting, under both training regimes. With training subjects only, the best classical model reached 0.561; adding the validation subjects raised it to 0.588, and refitting the GCN on the same subjects raised P3 × spatial to 0.613 (mean gain over all cells +0.005). The paired difference between each representation's best GCN cell and its best non-graph twin ranged from −0.044 to +0.050 in the train-only regime and from −0.049 to +0.033 in the train + val regime, and none was significant under either Holm-corrected test (Table S4; for P3, train only: +0.050, Wilcoxon–Holm $p = 0.401$; train + val: +0.033, $p = 0.801$). The GCN on wavelet features was numerically the best model in both regimes, but its advantage over non-graph classifiers is within the noise of this cohort. Nor could any of these comparisons be shown equivalent: with the corrected standard error the smallest difference detectable with 80% power was about 0.18 macro-F1 (Table S12), so an advantage of the graph model of that order can neither be shown nor excluded with 88 subjects. The spectral, wavelet, fused and raw representations cleared the demographic floor of 0.411; the connectivity (P2), time-domain (P4) and PSWE (P7) representations scored within about 0.05 of it.

Standard end-to-end EEG networks did no better. On the same raw windows and splits, EEGNet reached 0.530 ± 0.114 (ROC-AUC 0.754) and the 1D-CNN 0.523 ± 0.113 (Table S3; Figure S7), slightly above the GCN given the same raw input (+0.026 and +0.020; Wilcoxon–Holm $p = 0.517$) and below the GCN on wavelet features (−0.056, $p = 0.255$; −0.063, $p = 0.015$). The latter difference was significant for the 1D-CNN under the Wilcoxon test but under neither comparison with the Nadeau–Bengio correction. Learning features end to end from raw signal therefore did not substitute for a good hand-crafted representation on this cohort, while using 0.34 GPU-hours against none for every other model. As for the classical twins, these comparisons are imprecise (detectable difference about 0.25), so they show only that the deep networks are not dramatically better. Training EEGNet with the dropout recommended for cross-subject use (0.25 instead of the shared 0.5) changed little: macro-F1 0.551 ± 0.116, +0.021 against the original EEGNet (Wilcoxon–Holm $p = 0.767$) and still −0.035 below the GCN on wavelet features ($p = 0.705$; Supplementary methods).

4.6 Ablations

Table 4 and Figure S15 report each ablation on the confirmation folds as a paired difference from the same cell's confirmation-grid run. No ablation reached significance under either Holm-corrected test (0 of 32 with Wilcoxon, 0 with Nadeau–Bengio), and every point estimate except one was negative or zero: no modification improved on the reference configuration. Removing the graph altogether (identity adjacency) changed macro-F1 by −0.030 [−0.063, +0.004] for the spatial cell and by −0.022 for the hybrid cell, and a degree-preserving random graph by −0.012 and −0.004. Any benefit of message passing is therefore small, and we cannot attribute it to a specific electrode topology rather than to smoothing over neighbours. Two changes had uncorrected t-intervals that excluded zero for the spatial cell: GAT attention (−0.046 [−0.086, −0.006]) and a fully connected graph (−0.032 [−0.065, −0.000]), the latter consistent with over-smoothing; neither interval accounts for the overlap between cross-validation training sets. Sparsity ($k = 3$ or 6), binarised edges, depth and loss weighting changed macro-F1 by at most 0.021. Regressing age and sex out of every feature cost −0.023 [−0.066, +0.021] for the spatial cell, leaving performance well above the demographic floor.

Table 4: Ablations on the confirmation folds: paired Δ macro-F1 (ablated minus the same cell's confirmation-grid run) with 95% t-interval over the 25 splits, and Nadeau–Bengio p with Holm correction within each cell. Selection-fold values: Table S5.

	P3 × spatial		P3 × hybrid	
Ablation	Δ [95% CI]	p_{Holm}	Δ [95% CI]	p_{Holm}
Age/sex regressed out	-0.023 [-0.066, +0.021]	1.00	-0.037 [-0.080, +0.005]	1.00
No graph (identity)	-0.030 [-0.063, +0.004]	1.00	-0.022 [-0.056, +0.012]	1.00
Random graph (degree-preserving)	-0.012 [-0.040, +0.016]	1.00	-0.004 [-0.021, +0.013]	1.00
Fully connected	-0.032 [-0.065, -0.000]	1.00	-0.012 [-0.038, +0.015]	1.00
$k = 3$	-0.008 [-0.033, +0.016]	1.00	-0.002 [-0.022, +0.019]	1.00
$k = 6$	-0.012 [-0.040, +0.017]	1.00	-0.021 [-0.046, +0.005]	1.00
Binarised edges	-0.003 [-0.023, +0.018]	1.00	-0.016 [-0.045, +0.014]	1.00
Depth 1	-0.015 [-0.053, +0.022]	1.00	-0.002 [-0.041, +0.038]	1.00
Depth 3	-0.001 [-0.038, +0.036]	1.00	+0.001 [-0.040, +0.043]	1.00
GAT instead of GCN	-0.046 [-0.086, -0.006]	1.00	-0.030 [-0.076, +0.015]	1.00
20 s windows	-0.015 [-0.044, +0.015]	1.00	-0.005 [-0.036, +0.025]	1.00
5 s windows	-0.033 [-0.071, +0.005]	1.00	-0.008 [-0.046, +0.030]	1.00
No ICA artefact removal	-0.029 [-0.065, +0.008]	1.00	-0.035 [-0.076, +0.005]	1.00
No subject-equal weighting	-0.016 [-0.049, +0.016]	1.00	-0.013 [-0.043, +0.018]	1.00
+ P7 PSWE features	-0.008 [-0.040, +0.024]	1.00	-0.010 [-0.044, +0.025]	1.00
PSWE co-occurrence edges	-0.035 [-0.072, +0.001]	1.00	-0.027 [-0.062, +0.008]	1.00

Shorter (5 s) or longer (20 s) windows changed macro-F1 by −0.033 and −0.015 for the spatial cell (−0.008 and −0.005 for the hybrid cell), so 10 s windows were never improved upon. Removing ICA artefact rejection cost −0.029 [−0.065, +0.008] and −0.035, the largest of the preprocessing changes but still not significant. None of the ablations could be shown equivalent either (median detectable difference 0.14): they exclude large effects of these design choices, not small ones. The selection-fold ablations (Table S5) led to the same conclusions.

A learned graph does no better. Letting the GCN learn its adjacency, starting from the spatial graph, changed macro-F1 by between −0.016 and +0.005 across the seven representations on the selection folds (P3: 0.586, difference +0.000) and by +0.002 for P3 on the confirmation folds (0.598). The ranking of representations was unchanged (Spearman $\rho = 0.89$ with the fixed spatial graph). Of the 8 comparisons, 4 were formally equivalent within ±0.05 (P1, P2, P6, P7); the other four (P3, P4 and P5, and P3 on the confirmation folds) had corrected 90% intervals that included zero but reached ±0.075, and were inconclusive rather than different (Table S12). A data-driven graph therefore did not recover information that the fixed constructions missed.

Stopping rule. Re-running the whole grid with patience 50 or without early stopping left every conclusion unchanged: P3 × spatial remained the best cell, cell rankings were preserved (Spearman $\rho = 0.95$ and 0.81) and no representation's score changed significantly (Supplementary results).

4.7 Leakage inflates results to the top of the published range

Re-running the two best cells on the confirmation folds with window-level splitting — identical data, features, graph, model and hyperparameters — raised subject-level macro-F1 from 0.596 to 0.983 (spatial) and from 0.588 to 0.984 (hybrid). Window-level macro-F1, the metric most often reported, rose to 0.839 and 0.837 (Figure 3). The selection folds gave almost identical values (0.586 to 0.988 for the spatial cell). The leaky protocol thus reproduces the 95–99% figures reported with segment-level splits on this benchmark. It does not explain every published subject-wise result, which range from 54% to 84% for the three-class problem (Section 5.5).

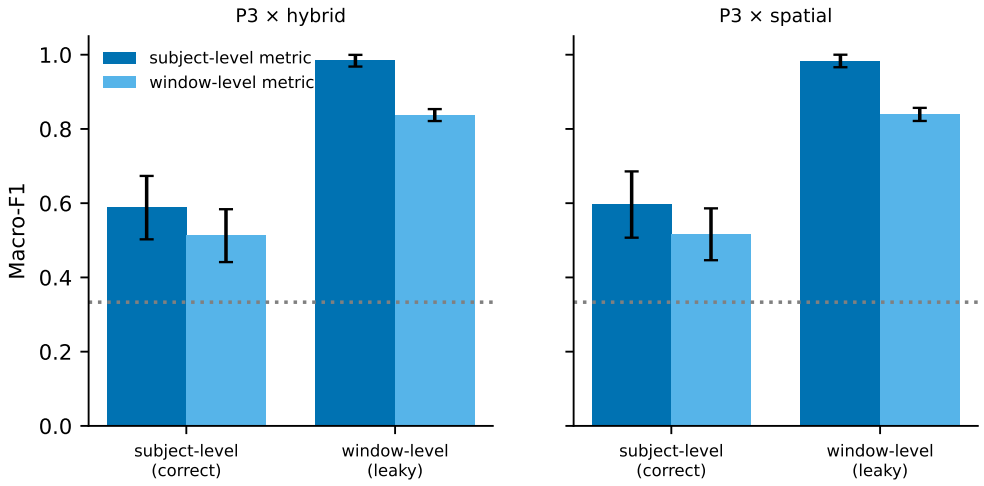


Figure 3: The same pipeline evaluated with subject-level (correct) and window-level (leaky) splitting on the confirmation folds. Bars: mean ± SD over 25 splits; dotted line, 1/3.

4.8 The best cell is above chance

On the first confirmation repeat, the observed macro-F1 of P3 × spatial was 0.624, against a label-permutation null with mean 0.312 and maximum 0.461 over 100 permutations ($p = 0.010$, the smallest value attainable; Figure S4). The original test on the selection folds gave the same result (200 permutations, observed 0.575, null maximum 0.474, $p = 0.005$).

4.9 PSWE: a BBB-associated marker that tracks slowing

Fixed-threshold PSWEs were most frequent in AD (median 1.34 events/min), intermediate in FTD (0.41) and least frequent in controls (0.22; Kruskal–Wallis $p = 0.041$; Figures S16 and 4). In negative-binomial models adjusted for age and sex, the PSWE rate ratio relative to controls was 2.96 (95% CI 1.58–5.57, $p < 0.001$) in AD and 2.07 (1.03–4.15, $p = 0.041$) in FTD.

The PSWE rate was, however, almost collinear with background slowing (Spearman $\rho = 0.88$ with relative delta+theta power; age- and sex-adjusted partial $\rho = 0.87$). Once slowing was added to the model, the AD rate ratio fell below one, to 0.50 (0.33–0.75, $p < 0.001$), and FTD to 0.67 ($p = 0.051$). This is not an artefact of multicollinearity among the predictors — the variance inflation of the slowing covariate given group, age and sex was only 1.31 — but it does rest on extrapolation: 50% of AD and 26% of FTD patients were more slowed than any control. Restricted to the 18 AD and 17 FTD patients whose slowing lay within the control range, no PSWE excess was detectable even before adjusting for slowing (rate ratios 0.62, 95% CI 0.30–1.28, $p = 0.196$, and 1.12, 0.54–2.33, $p = 0.770$); these intervals are wide and do not exclude a moderate excess. Relative to the events that their slowing predicts — the slowing–PSWE relation estimated in controls only — patients had fewer events than controls (median observed/expected 0.45 for AD and 0.58 for FTD versus 0.86; Kruskal–Wallis $p < 0.001$). The PSWE rate was not significantly associated with MMSE within patients ($\rho = -0.22$, $p = 0.095$). In the classifier, PSWE features alone (P7) performed near the demographic floor (0.448), and adding them to the best representation, or using PSWE co-occurrence as the graph, did not help (Section 4.6).

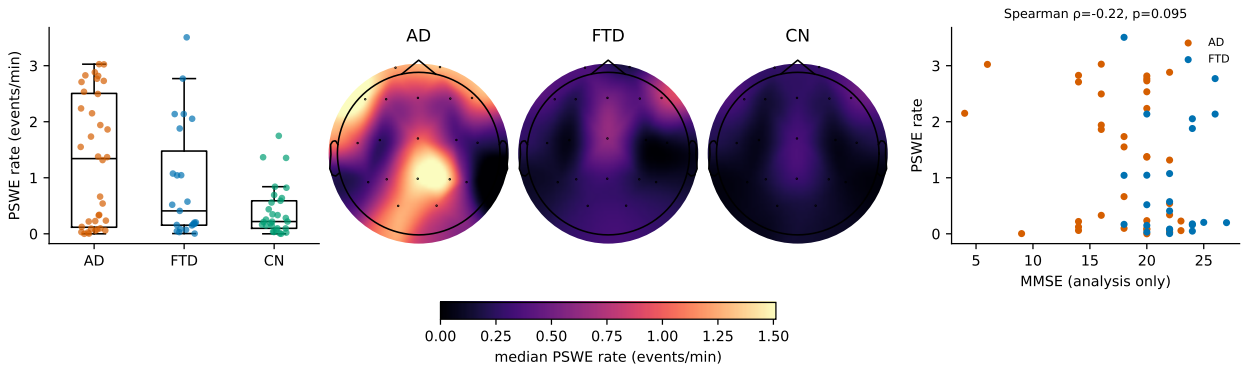


Figure 4: Fixed-threshold PSWE rate per subject by group (left), median per-channel rate as scalp maps (middle), and rate against MMSE in patients (right; analysis only).

Detector robustness. Neither the coarse MPF grid nor artefacts explain these results. With zero-padded spectra and artefact rejection together (a median of at most 0.3% of segments rejected per group), rates correlated $\rho = 0.98$ with the original ones, the age- and sex-adjusted rate ratios were 4.54 for AD ($p < 0.001$) and 2.85 for FTD ($p = 0.008$), and adjusting for slowing again brought them below one (0.50 and 0.62). Without ICA, the rate ratios were 2.35 and 1.99 ($\rho = 0.90$ with slowing), so ocular artefacts do not create the excess either. Nor does the event rule (Table S13): across the 9 combinations of 5–7 Hz thresholds and 3–8 s minimum durations (A11), the age- and sex-adjusted AD rate ratio ranged from 2.14 to 3.95 (significant in 9 of 9), FTD from 1.37 to 2.75 (6 of 9), the correlation with slowing from 0.80 to 0.91, and every slowing-adjusted AD ratio was below one (0.24–0.78).

Background-relative PSWEs (exploratory). Defined relative to each channel's own background, transient slowing events were less frequent in AD and FTD than in controls under every variant (Supplementary results).

4.10 External validation: Parkinson's disease (ds004584)

The ds004584 recordings passed the same quality gates (2,327 windows from 149 subjects; 2.2 ICA components removed per recording on average), and theta/alpha slowing was again higher in patients (0.95) than in controls (0.72). Age and sex alone reached a macro-F1 of 0.550 against 0.490 for chance in this two-class problem.

Protocol effects replicated. Here the representation effect survived the primary, overlap-adjusted test ($F_{\mathrm{NB}} = 2.63$, $p = 0.031$), unlike on ds004504. Descriptively, in the by-subject mixed model, $F(6, 882) = 20.03$ with 13 of 21 pairs separated, whereas graph construction was small (marginal means 0.587 spatial, 0.593 functional, 0.590 hybrid; $p = 0.002$) and favoured a different graph type than on ds004504; on macro-F1 with splits as units, representation was significant ($F = 23.19$, $p < 0.001$, partial $\eta^2 = 0.49$) and graph construction was not ($p = 0.207$) in the unadjusted analysis. The GCN was not significantly better than non-graph classifiers trained on the same subjects (differences -0.008 to $+0.049$, none significant after Holm correction). Window-level splitting of the configuration selected on ds004504 inflated subject-level macro-F1 from 0.688 to 0.839 (window-level 0.810), in line with an independent audit of the same cohort [27].

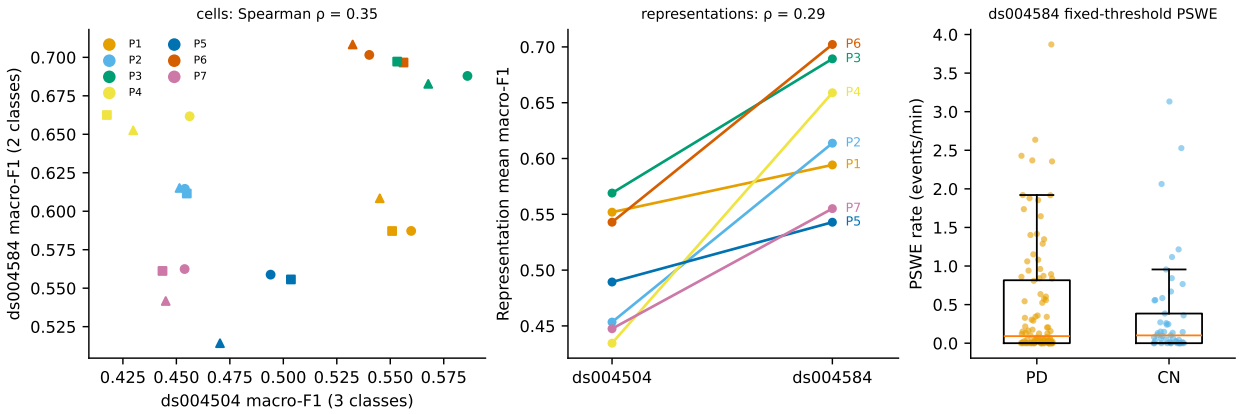


Figure 5: External validation on ds004584. Left: every cell's mean macro-F1 on ds004504 against ds004584. Middle: representation means on both datasets. Right: fixed-threshold PSWE rate per subject in Parkinson's disease and controls.

The ranking only partly transferred. The configuration selected on ds004504 (P3 × spatial), evaluated without re-selection, reached macro-F1 0.688 ± 0.101 (subject-bootstrap 95% CI 0.651–0.721; balanced accuracy 0.709, ROC-AUC 0.766), well above its permutation null (observed 0.679 on the first repeat against a null maximum of 0.593; $p = 0.010$, 100 permutations) and ranked 5th of 21; the best cell was P6 × hybrid (0.708). Wavelet and fused features were among the top three on both datasets (Table S11), but the full ordering did not transfer: Spearman $\rho = 0.35$ ($p = 0.122$) across the 21 cells and $\rho = 0.29$ ($p = 0.535$) across the seven representations, mainly because time-domain (P4) and connectivity (P2) features, near the bottom on ds004504, ranked third and fourth on ds004584, and raw windows (P5) fell to last (Figure 5).

PSWEs in Parkinson's disease. Fixed-threshold PSWEs were not more frequent in Parkinson's disease than in controls (median 0.09 versus 0.10 events/min; age- and sex-adjusted rate ratio 1.21, 95% CI 0.65–2.23, $p = 0.552$), consistent with the earlier report that they are uncommon in this disease [34]. As on ds004504, the rate tracked background slowing ($\rho = 0.81$), and adjusting for slowing brought the ratio below one (0.57, $p = 0.005$). The rate was not associated with MoCA within patients ($\rho = -0.11$, $p = 0.297$), no background-relative variant differed between groups (smallest $p = 0.186$), and zero-padding with artefact rejection left the result unchanged (rate ratio 1.17, $p = 0.674$). No event rule of the A11 grid produced a significant Parkinson's excess (rate ratios 0.89–1.75), almost all patients lay within the controls' range of slowing (95 of 100), and relative to their slowing they had fewer events than controls (median observed/expected 0.42 versus 0.73, $p = 0.007$).

5 Discussion

5.1 Principal findings

Holding the model, data and evaluation protocol fixed, the factors we varied differ widely in how far they move macro-F1: splitting windows instead of subjects by 0.39, the choice of representation by up to 0.13 (averaged over graph types), and the choice of graph by up to 0.015. The title's 'not' therefore means 'far less than', not 'zero'. In detail we found four things. First, in the primary, overlap-adjusted test the representation effect was not significant on ds004504 ($p = 0.258$ and

$p = 0.188$) but was in the external cohort ($p = 0.031$). Descriptively, representation had the largest effect on three-class AD/FTD/control classification (by-subject $F = 16.60$ and 19.81; split-level partial $\eta^2 = 0.33$ and 0.36), with wavelet, spectral and fused features leading, but in simulation these unadjusted tests declared a difference between equally informative representations in 34–80% of datasets. A representation effect is therefore established in Parkinson's disease, not in this dementia cohort, where 88 subjects cannot distinguish the observed spread from sampling variation (a null simulation matched to the data puts it at $p = 0.31$). Second, how the graph is built had a negligible effect: with subjects as units it was descriptively significant but spans at most 0.007 in the probability assigned to the true class, it changes direction between cohorts and does not survive the adjustment for overlapping training sets, averaged over representations the three graph constructions were formally equivalent within ±0.05 macro-F1, and within single representations the comparisons were inconclusive rather than different. The GCN showed no detectable advantage over a random forest or SVM trained on the same features and subjects, and end-to-end deep EEG networks (EEGNet, a 1D-CNN) on raw windows did not reach it either; these comparisons, however, could only have detected differences of about 0.18–0.25 macro-F1, so moderate advantages in either direction remain possible. These conclusions held under three stopping rules, with 5 s to 20 s windows and without ICA. Third, splitting windows instead of subjects inflates subject-level macro-F1 from 0.596 to 0.983 on the identical pipeline, reproducing the top of the published range, although not every published subject-wise figure (Section 5.5). Fourth, the blood–brain-barrier-associated PSWE marker is more frequent in AD and FTD than in controls, but this excess is largely accounted for by continuous background slowing; defined relative to each person's background, transient slowing events are, if anything, less frequent in dementia.

5.2 Representation and topology

On effect sizes, representation outweighed graph construction in every analysis (partial η^2 0.33 against 0.14; adjusted F 1.34 against 0.46), and the external cohort confirms that representation can matter; on ds004504 neither effect is statistically established. Within that limit, the information a GNN can use on this benchmark lives mostly in the node features. Several observations point the same way. Spatial graphs — which are identical for every subject and therefore carry no diagnostic information of their own — did at least as well as functional and hybrid graphs, and averaged over representations the three graph types were formally equivalent; within single representations the design could only have detected differences of about 0.12 (none of 42 comparisons was shown equivalent or different). Neither removing the graph nor replacing it with a degree-preserving random graph changed performance significantly; any benefit of message passing is small, and we cannot attribute it to a particular connectivity pattern rather than to smoothing over neighbouring electrodes. And non-graph classifiers on flattened features, trained on the same subjects, were not significantly different from the GCN for any representation, although the GCN on wavelet features was numerically best in both training regimes and these comparisons cannot exclude differences of about 0.18. These results agree with the independent conclusion of the 2026 scoping review that apparent representational gains on this dataset do not survive leakage-free evaluation [22], and they temper the case for graph-based modelling made in recent work [21, 28, 29]: on 19-electrode scalp EEG from 88 subjects and with this compact, untuned GCN, we see no evidence that graph structure adds information beyond good node features. A graph learned from the data, starting from the spatial one, did no better either (Section 4.6). These conclusions are scoped to the architecture we fixed: a two-layer GCN of about 3k parameters, with fixed hyperparameters, on 19 scalp electrodes, and to the graphs we built: one functional graph (alpha-band wPLI), its product with the spatial graph, and one learned adjacency. Whether they hold for tuned or deeper graph models, functional

graphs from other bands (theta mattered most to the best model), attention over time, source-space nodes or much larger cohorts remains open.

5.3 Neurophysiological interpretation

Representations that describe each electrode's own spectrum (P1, P3) scored higher than one built only from inter-electrode phase coupling (P2), and the functional graph did not improve on the spatial one. These differences are descriptive — none survived the adjustment for overlapping training sets — so they are consistent with, rather than evidence for, local oscillatory slowing being a stronger discriminative signal than altered large-scale coupling, a reading in line with spectral slowing being the most replicated EEG marker of AD [39, 40] and with the posterior-alpha-to-diffuse-theta shift visible in our data (Figure S10). Occlusion of the best model points the same way: it relied mainly on the 4–8 and 8–16 Hz wavelet bands, recognised controls chiefly through theta, and depended on no single electrode (Figure S8). It does not contradict network accounts of dementia [6, 7]: alpha-band wPLI estimated within 10 s windows is noisy, and electrode-level graphs are a coarse proxy for cortical networks. The error pattern is clinically recognisable: controls are identified best, while FTD — the smallest group, and less slowed than AD — is identified worst and is usually mistaken for AD. Separating the two dementias remains the hard part of the problem.

5.4 What the PSWE results mean for an EEG marker of BBB dysfunction

With the fixed threshold, our data agree in direction with Milikovsky et al.: PSWE rates were roughly three times higher in AD than in controls after adjusting for age and sex, and FTD — not previously studied — was also elevated. Two further analyses change the interpretation. The fixed-threshold rate was almost collinear with relative delta+theta power; half of the AD patients were more slowed than any control, and among patients whose slowing overlapped with controls no excess was detectable (AD rate ratio 0.62, 95% CI 0.30–1.28), while relative to the events their slowing predicts, patients had fewer. When events were instead defined as transient drops below each channel's own background, AD and FTD showed fewer of them than controls, under every variant we examined. Together these results suggest that, in resting scalp EEG with diffuse slowing, a median power frequency below 6 Hz mostly flags a slow *background* rather than *paroxysmal* events.

We draw three cautious conclusions. First, a fixed-threshold PSWE rate should not be read as evidence of BBB dysfunction without controlling for background slowing. Second, background-relative definitions appear necessary if PSWEs are to capture transient cortical dysfunction, although they bring their own floor effect in very slow recordings, which we observed in the Hz-drop variants. Third, our data cannot test the BBB link itself: ds004504 contains no BBB imaging, the original detection details include parameters we could not verify from the full text, and Milikovsky et al. studied different patients under different recording conditions. The decisive test is a cohort with both EEG and dynamic contrast-enhanced MRI.

5.5 How results on this benchmark should be read

On one pipeline, window-level instead of subject-level splitting moved subject-level macro-F1 from 0.596 to 0.983, far more than any representation, graph or architecture change we tested. Figures in the high 90s on ds004504 should therefore be read alongside their splitting protocol, as argued for this benchmark [22], for translational EEG [26, 27, 47] and for machine-learning-based science generally [23–25].

Leakage does not, however, account for every published figure. Table 5 lists the subject-wise results we could verify from the full texts. Our three-class result (macro-F1 0.596, accuracy 0.623)

Table 5: Published subject-wise results on the AHEPA AD/FTD/control EEG cohort, verified against each paper's full text, and this study. Accuracy unless stated. LNSO: leave-N-subjects-out; LOSO: leave-one-subject-out. Results are not directly comparable: tasks, scored units and validation differ.

Study	Task	Validation	Scored unit	Result
Miltiadous [9]	AD/CN; FTD/CN	LOSO	–	83.3%; 75.0%
Kang [65]	AD/CN	LNSO	–	74.2%
Stefanou [20]	AD/FTD/CN	LNSO (five subsets per group)	–	54.3% (macro-F1 0.49)
Chen [64]	AD/FTD/CN	10-fold, split unit not stated; decision combined with MMSE	subject	80.2%
Yıldız [21]	AD/FTD/CN	subject-wise 10-fold; ds006036	20-s segment (50% overlap)	84.3% (macro-F1 0.84)
This study	AD/FTD/CN	subject-level 5×5 CV, confirmation folds	subject	macro-F1 0.596; accuracy 0.623

is close to the 54% that the dataset's authors report under leave-subjects-out validation [20], and well below the 80–84% of two other reports [21, 64]. Those differ from our protocol in ways that could each matter. One combines the EEG model's output with the MMSE [64], on which every control in this cohort scores 30. The other scores overlapping 20 s segments rather than subjects, uses OpenNeuro ds006036 rather than ds004504, and evaluates subject-wise a configuration chosen among 15 in earlier experiments [21]; per-segment accuracy is not comparable to our per-subject macro-F1, although our own per-segment accuracy under subject-level splits (0.544) is lower, not higher, than our per-subject accuracy. From the published descriptions we cannot tell whether preprocessing or normalisation fitted across subjects also contributes. Each of these differences — what is scored, what enters the decision, which recordings are used, how the configuration was chosen — is itself a choice of evaluation protocol, so the spread of subject-wise figures is further evidence for the paper's title rather than an exception to it. Our design compares conditions under one compact, untuned model; it is not a search for the best classifier, and a tuned or pretrained model could well score higher under the same protocol. At the level reached here (FTD sensitivity 0.368, expected calibration error 0.265), EEG classifiers of this kind are research benchmarks, not diagnostic tools. The permutation test and the confirmation folds show only that the result is real and not an artefact of choosing the best of 21 configurations.

5.6 What the external cohort adds

Applying the identical protocol to Parkinson's disease, a different site, device and recording state, reproduced the protocol-level findings — a representation effect (established there after the adjustment for overlapping training sets), a negligible graph effect, no advantage of the GCN over non-graph twins, inflation under window-level splitting, and a fixed-threshold PSWE rate that tracks slowing — but not the detailed ranking of representations ($\rho = 0.29$). The configuration chosen on dementia generalised reasonably (0.688, 5th of 21), and wavelet and fused features were near the top in both cohorts, but time-domain and connectivity features, weak for dementia, were competitive for Parkinson's disease. This argues against treating any representation ranking from one disease as a general recommendation: the evaluation protocol transfers; the best representation is disease- and

possibly recording-specific.

5.7 Limitations

The cohort is small, and cross-validation estimates in samples of this size carry wide error bars [49]: test folds contain 17–18 subjects, and under the conservative Nadeau–Bengio correction no pairwise difference between representations, and no single ablation, reached significance, and the omnibus representation effect on ds004504 did not survive the same adjustment; equivalence tests show that, apart from the averaged graph-type comparisons, the design could only have detected differences of roughly 0.14–0.18 macro-F1 between individual configurations. We report both corrected and uncorrected tests and give effect sizes, but small differences between leading representations should not be over-interpreted. All data come from one site, device and protocol, and FTD is the smallest class ($n = 23$). Diagnoses are clinical; the dataset reports neither medication (cholinesterase inhibitors and psychotropic drugs alter the EEG), nor FTD subtype, nor biomarker confirmation, so label noise and treatment effects cannot be assessed. The external cohort addresses generalisation of the protocol but differs in disease, task (eyes open), mains frequency and recording length (median about 2.5 min), was recorded with patients on their usual levodopa [60], so medication and disease effects on the EEG cannot be separated, and is two-class, so it cannot test whether the AD/FTD results themselves replicate. The confirmation folds re-sample the same 88 subjects, so they remove the selection bias of choosing among configurations but are not an external replication, and no independent AD/FTD cohort recorded with a comparable protocol was available to us. Our conclusions concern one compact model family: the learned-adjacency variant kept the representation ranking ($\rho = 0.89$), but the GAT check covered only the two reference cells, and larger or pretrained models might separate representations that this GCN does not. Hyperparameters were fixed a priori and not tuned. Early stopping on a validation set of about 14 subjects selected a median of 9 epochs (23% of runs stopped at epoch 3 or earlier; median 5 for P1 and 20 for P3), so many models are lightly trained; however, training longer (patience 50, or 100 epochs without early stopping) changed no conclusion and did not help the highest-capacity configuration (raw windows, about 84k parameters), and refitting on training plus validation subjects for the selected epoch count changed macro-F1 by only +0.005 on average. The adjusted test is conservative (no false positives in simulation, power of 20–49% for moderate differences), so 'not established' is weak evidence that representation does not matter; the out-of-bootstrap comparison, which does not rely on it, reached the same conclusion, with a bootstrap standard error of about 0.095 for a paired difference, so that only differences of roughly twice that size could be shown with this cohort. Calibration was poor. Results are reproduced exactly from the original feature caches (Kaggle notebook output, too large for the repository), but rebuilding the caches from the raw recordings in a later Kaggle environment changed per-split macro-F1 by 0.065 on average (mean 0.571 against 0.586 for the best cell; 87% of subject decisions unchanged), a reminder that preprocessing libraries are part of the pipeline. Graphs are defined on electrodes rather than cortical sources. The demographic floor (0.411) shows residual age and sex information, which residualisation reduced but did not remove. For PSWEs, the MPF frequency range was not reported in the source and was set by us (zero-padding the spectrum and rejecting artefacts did not change the results), we detected events per channel rather than over averaged regions, the background-relative analysis is post hoc, and no BBB measure is available.

5.8 Future work

The most valuable next study is a multi-site replication of this exact protocol, since a single-site design cannot establish device or population invariance. Natural extensions are source-space graphs

[29], richer learned-graph models than our single trainable adjacency [13], spatio-temporal graphs over window sequences [14, 16], an architecture sweep under a fixed representation (the transpose of this design) [11], an MCI class, and — for the BBB question — simultaneous EEG and dynamic contrast-enhanced MRI to test whether any PSWE definition tracks measured barrier leakage [30, 31].

6 Conclusion

We asked whether the performance of graph neural networks for EEG-based dementia classification comes from how the signal is represented, how the brain graph is built, or the model. On a public 88-subject benchmark, with a compact, untuned model, the data and the evaluation protocol held fixed and seven representations crossed with three graph constructions over 25 paired subject-level splits, the representation had by far the larger effect: wavelet and spectral node features led, graph construction had a detectable but negligible effect, and the GCN did not outperform non-graph classifiers on the same features. Once the overlap between training sets is accounted for — uncorrected tests, a simulation showed, find differences between equally informative representations in most datasets — the representation effect is established in an external Parkinson's disease cohort but not on ds004504, and which representation works best did not transfer between diseases. The best configuration, chosen on one set of splits and re-evaluated on an independent set, reached a three-class macro-F1 of 0.596, clearly above chance ($p = 0.010$). Window-level splitting inflated it to 0.983 on the same pipeline, reproducing the upper end of published figures; published subject-wise results range more widely than leakage alone explains. A blood–brain-barrier-associated EEG marker, the paroxysmal slow-wave event, was elevated in AD and FTD under its original definition, but the excess tracked background slowing and was not detectable among patients whose slowing overlapped with controls'; it was not elevated in Parkinson's disease, and relative to each person's background, transient slowing events were less frequent in disease. We release the code, seeds, frozen folds and every raw run so that each number in this paper can be regenerated from the data.

Data and code availability

The EEG data are OpenNeuro ds004504 (CC0) [19] and ds004584 [59]. Code, configuration (`configs/default.yaml`), the frozen fold assignment (`splits/folds_ds004504.json`, with the SHA-256 of `participants.tsv`), every raw training run (one JSON record per run, with configuration hash, code version, seed and per-subject test predictions), and the Kaggle notebooks that produced them are available at `https://github.com/Devguru-codes/Major_project_IIITN`. Random seeds are 0–4 for the main grid (selection folds), 5–9 for the confirmation folds (`splits/folds_ds004504_confirm.json`, committed before any confirmation run), and 10,000 + i and 11,000 + i for permutation i on the selection and confirmation folds. The main tables and figures are regenerated from the committed raw runs by `python -m eegrep report`; the later analyses by the Kaggle notebooks NB10–NB23 (each a sequence of `eegrep` commands listed in `kaggle/build_notebooks.py`); and every number in the text by `paper/make_numbers.py`, which reads only committed outputs. The resolved package versions are recorded in `reports/kaggle_pip_freeze.txt`. The preprocessed feature caches (output of NB02) are not in the repository because of their size; rebuilding them from the raw recordings in a later software environment changes results slightly (Section 5.7). The full study runs on free cloud CPUs, except the deep-learning baselines (0.57 GPU-hours).

Ethics statement

This study is a secondary analysis of two public, de-identified datasets. Recording of ds004504 was approved by the Scientific and Ethics Committee of AHEPA University Hospital, Aristotle University of Thessaloniki (protocol 142/12-04-2023), with informed consent from all participants [18]. Recording of ds004584 was approved by the University of Iowa Human Subjects Review Board (IRB 201707828), with written informed consent from all participants [60]. [**Author to complete: state whether your institution required further review for secondary analysis of public de-identified data.**]

Acknowledgements

We thank the authors of ds004504 and ds004584 for releasing their recordings openly, and OpenNeuro for hosting them. All computation used free Kaggle notebook sessions.

Funding

[**Author to complete.**]

Competing interests

[**Author to complete.**]

Author contributions (CRediT)

[**Author to complete: conceptualisation, methodology, software, formal analysis, writing, supervision.**]

Use of generative AI

[**Author to complete: most journals require disclosure of generative-AI assistance in code development, analysis and manuscript preparation, and the authors remain responsible for all content.**]

References

[1] GBD 2019 Dementia Forecasting Collaborators. Estimation of the global prevalence of dementia in 2019 and forecasted prevalence in 2050: an analysis for the Global Burden of Disease Study 2019. *The Lancet Public Health*, 7(2):e105–e125, 2022. doi: 10.1016/S2468-2667(21)00249-8. URL `https://doi.org/10.1016/S2468-2667(21)00249-8`.

[2] Alzheimer's Association. 2024 Alzheimer's disease facts and figures. *Alzheimer's & Dementia*, 20(5):3708–3821, 2024. doi: 10.1002/alz.13809. URL `https://doi.org/10.1002/alz.13809`.

[3] David B. Hogan, Nathalie Jett'e, Kirsten M. Fiest, Jodie I. Roberts, Dawn Pearson, Eric E. Smith, Pamela Roach, Andrew Kirk, Tamara Pringsheim, and Colleen J. Maxwell. The prevalence and incidence of frontotemporal dementia: a systematic review. *Canadian Journal*

of Neurological Sciences, 43(S1):S96–S109, 2016. doi: 10.1017/cjn.2016.25. URL https://doi.org/10.1017/cjn.2016.25.

[4] Stevie Hendriks, Kirsten Peetoom, Christian Bakker, Wiesje M. van der Flier, Janne M. Papma, Raymond Koopmans, Frans R. J. Verhey, Marjolein de Vugt, Sebastian Köhler, Adrienne Withall, Juliette L. Parlevliet, Özgül Uysal-Bozkir, Roger C. Gibson, Susanne M. Neita, Thomas Rune Nielsen, Lise C. Salem, Jenny Nyberg, Marcos Antonio Lopes, Jacqueline C. Dominguez, Ma Fe De Guzman, Alexander Egeberg, Kylie Radford, Tony Broe, Mythily Subramaniam, Edimansyah Abdin, Amalia C. Bruni, Raffaele Di Lorenzo, Kate Smith, Leon Flicker, Merel O. Mol, Maria Basta, Doris Yu, Golden Masika, Maria S. Petersen, and Luis Ruano. Global prevalence of young-onset dementia: A systematic review and meta-analysis. *JAMA Neurology*, 78(9):1080–1090, 2021. doi: 10.1001/jamaneurol.2021.2161. URL https://doi.org/10.1001/jamaneurol.2021.2161.

[5] J. Cummings, L. Apostolova, G.D. Rabinovici, A. Atri, P. Aisen, S. Greenberg, S. Hendrix, D. Selkoe, M. Weiner, R.C. Petersen, and S. Salloway. Lecanemab: Appropriate use recommendations. *The Journal of Prevention of Alzheimer's Disease*, 10(3):362–377, 2023. doi: 10.14283/jpad.2023.30. URL https://doi.org/10.14283/jpad.2023.30.

[6] C. Stam, B. Jones, G Nolte, M Breakspear, and P. Scheltens. Small-world networks and functional connectivity in Alzheimer's disease. *Cerebral Cortex*, 17(1):92–99, 2007. ISSN 1460-2199. doi: 10.1093/cercor/bhj127. URL http://dx.doi.org/10.1093/cercor/bhj127.

[7] Willem de Haan, Yolande AL Pijnenburg, Rob LM Strijers, Yolande van der Made, Wiesje M van der Flier, Philip Scheltens, and Cornelis J Stam. Functional neural network analysis in frontotemporal dementia and Alzheimer's disease using EEG and graph theory. *BMC Neuroscience*, 10(1), 2009. ISSN 1471-2202. doi: 10.1186/1471-2202-10-101. URL http://dx.doi.org/10.1186/1471-2202-10-101.

[8] Cosimo Ieracitano, Nadia Mammone, Amir Hussain, and Francesco C. Morabito. A novel multi-modal machine learning based approach for automatic classification of EEG recordings in dementia. *Neural Networks*, 123:176–190, 2020. ISSN 0893-6080. doi: 10.1016/j.neunet.2019.12.006. URL http://dx.doi.org/10.1016/j.neunet.2019.12.006.

[9] Andreas Miltiadous, Emmanouil Gionanidis, Katerina D. Tzimourta, Nikolaos Giannakeas, and Alexandros T. Tzallas. DICE-Net: A novel convolution-transformer architecture for Alzheimer detection in EEG signals. *IEEE Access*, 11:71840–71858, 2023. ISSN 2169-3536. doi: 10.1109/access.2023.3294618. URL http://dx.doi.org/10.1109/access.2023.3294618.

[10] Rundong Jiang, Xiaowei Zheng, Jiamin Sun, Lei Chen, Guanghua Xu, and Rui Zhang. Classification for Alzheimer's disease and frontotemporal dementia via resting-state electroencephalography-based coherence and convolutional neural network. *Cognitive Neurodynamics*, 19(1), 2025. ISSN 1871-4099. doi: 10.1007/s11571-025-10232-2. URL http://dx.doi.org/10.1007/s11571-025-10232-2.

[11] Dominik Klepl, Min Wu, and Fei He. Graph neural network-based EEG classification: A survey. *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, 32:493–503, 2024. ISSN 1558-0210. doi: 10.1109/tnsre.2024.3355750. URL http://dx.doi.org/10.1109/tnsre.2024.3355750.

[12] Dominik Klepl, Fei He, Min Wu, Daniel J. Blackburn, and Ptolemaios Sarrigiannis. EEG-based graph neural network classification of Alzheimer's disease: An empirical evaluation of functional connectivity methods. *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, 30:2651–2660, 2022. ISSN 1558-0210. doi: 10.1109/tnsre.2022.3204913. URL http://dx.doi.org/10.1109/tnsre.2022.3204913.

[13] Dominik Klepl, Fei He, Min Wu, Daniel J. Blackburn, and Ptolemaios Sarrigiannis. Adaptive gated graph convolutional network for explainable diagnosis of Alzheimer's disease using EEG data. *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, 31:3978–3987, 2023. ISSN 1558-0210. doi: 10.1109/tnsre.2023.3321634. URL http://dx.doi.org/10.1109/tnsre.2023.3321634.

[14] Xiaocai Shan, Jun Cao, Shoudong Huo, Liangyu Chen, Ptolemaios Georgios Sarrigiannis, and Yifan Zhao. <scp>spatial–temporal</scp> graph convolutional network for Alzheimer classification based on brain functional connectivity imaging of electroencephalogram. *Human Brain Mapping*, 43(17):5194–5209, 2022. ISSN 1097-0193. doi: 10.1002/hbm.25994. URL http://dx.doi.org/10.1002/hbm.25994.

[15] Qingjie Xu, Libing An, Haiqiang Yang, and Keum-Shik Hong. A multi-graph convolutional network method for Alzheimer's disease diagnosis based on multi-frequency EEG data with dual-mode connectivity. *Frontiers in Neuroscience*, 19, 2025. ISSN 1662-453X. doi: 10.3389/fnins.2025.1555657. URL http://dx.doi.org/10.3389/fnins.2025.1555657.

[16] Ehssan Aljanabi and İlker Türker. Connectogram-coh: A coherence-based time-graph representation for EEG-based Alzheimer's disease detection. *Diagnostics*, 15(11):1441, 2025. ISSN 2075-4418. doi: 10.3390/diagnostics15111441. URL http://dx.doi.org/10.3390/diagnostics15111441.

[17] Yujian Liu, Libing An, Haiqiang Yang, and Shuzhi Sam Ge. Multi-frequency EEG and multi-functional connectivity graph convolutional network based detection method of patients with Alzheimer's disease. *Complex & Intelligent Systems*, 11(8), 2025. ISSN 2198-6053. doi: 10.1007/s40747-025-01974-x. URL http://dx.doi.org/10.1007/s40747-025-01974-x.

[18] Andreas Miltiadous, Katerina D. Tzimourta, Theodora Afrantou, Panagiotis Ioannidis, Nikolaos Grigoriadis, Dimitrios G. Tsalikakis, Pantelis Angelidis, Markos G. Tsipouras, Euripidis Glavas, Nikolaos Giannakeas, and Alexandros T. Tzallas. A dataset of scalp EEG recordings of Alzheimer's disease, frontotemporal dementia and healthy subjects from routine EEG. *Data*, 8 (6):95, 2023. ISSN 2306-5729. doi: 10.3390/data8060095. URL http://dx.doi.org/10.3390/data8060095.

[19] Andreas Miltiadous, Katerina D. Tzimourta, Theodora Afrantou, Panagiotis Ioannidis, Nikolaos Grigoriadis, Dimitrios G. Tsalikakis, Pantelis Angelidis, Markos G. Tsipouras, Evripidis Glavas, Nikolaos Giannakeas, and Alexandros T. Tzallas. A dataset of EEG recordings from: Alzheimer's disease, frontotemporal dementia and healthy subjects, 2024. URL https://openneuro.org/datasets/ds004504/versions/1.0.8.

[20] Konstantinos Stefanou, Katerina D. Tzimourta, Christos Bellos, Georgios Stergios, Konstantinos Markoglou, Emmanouil Gionanidis, Markos G. Tsipouras, Nikolaos Giannakeas, Alexandros T. Tzallas, and Andreas Miltiadous. A novel CNN-based framework for Alzheimer's disease detection using EEG spectrogram representations. *Journal of Personalized Medicine*, 15(1):27, 2025. doi: 10.3390/jpm15010027. URL https://doi.org/10.3390/jpm15010027.

[21] Abdulnasır Yıldız and Hasan Zan. Deep learning-based Alzheimer's disease detection from multi-channel EEG using fused time–frequency image grids. *Diagnostics*, 16(5):746, 2026. ISSN 2075-4418. doi: 10.3390/diagnostics16050746. URL http://dx.doi.org/10.3390/diagnostics16050746.

[22] Andreas Miltiadous, Aimilia Ntetska, Vasileios Aspiotis, Efthalia Moustakli, Markos G. Tsipouras, Alexandros T. Tzallas, Nikolaos Giannakeas, Euripidis Glavas, Pantelis Angelidis, and Katerina D. Tzimourta. The AHEPA EEG benchmark: setting the standard for machine learning in dementia diagnosis, a scoping review. *Cognitive Neurodynamics*, 20(1), 2026. ISSN 1871-4099. doi: 10.1007/s11571-026-10464-w. URL http://dx.doi.org/10.1007/s11571-026-10464-w.

[23] Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns*, 4(9):100804, 2023. ISSN 2666-3899. doi: 10.1016/j.patter.2023.100804. URL http://dx.doi.org/10.1016/j.patter.2023.100804.

[24] Matthew Rosenblatt, Link Tejavibulya, Rongtao Jiang, Stephanie Noble, and Dustin Scheinost. Data leakage inflates prediction performance in connectome-based machine learning models. *Nature Communications*, 15(1):1829, 2024. doi: 10.1038/s41467-024-46150-w. URL https://doi.org/10.1038/s41467-024-46150-w.

[25] Sohrab Saeb, Luca Lonini, Arun Jayaraman, David C. Mohr, and Konrad P. Kording. The need to approximate the use-case in clinical machine learning. *GigaScience*, 6(5):gix019, 2017. doi: 10.1093/gigascience/gix019. URL https://doi.org/10.1093/gigascience/gix019.

[26] Geoffrey Brookshire, Jake Kasper, Nicholas M. Blauch, Yunan Charles Wu, Ryan Glatt, David A. Merrill, Spencer Gerrol, Keith J. Yoder, Colin Quirk, and Ché Lucero. Data leakage in deep learning studies of translational EEG. *Frontiers in Neuroscience*, 18:1373515, 2024. doi: 10.3389/fnins.2024.1373515. URL https://doi.org/10.3389/fnins.2024.1373515.

[27] Burak Aggul. Window leakage in clinical EEG machine learning: matched subject-independent evaluation across ADHD and Parkinson cohorts. *Frontiers in Neuroinformatics*, 20:1930795, 2026. doi: 10.3389/fninf.2026.1930795. URL https://doi.org/10.3389/fninf.2026.1930795.

[28] Thawirasm Jungrungrueang, Sawrawit Chairat, Kasidach Rasitanon, Praopim Limsakul, and Krit Charupanit. Translational approach for dementia subtype classification using convolutional neural network based on EEG connectome dynamics. *Scientific Reports*, 15(1), 2025. ISSN 2045-2322. doi: 10.1038/s41598-025-02018-7. URL http://dx.doi.org/10.1038/s41598-025-02018-7.

[29] Shuzhi Zhao, Guangyan Dai, Jingting Li, Xiaoxia Zhu, Xiyan Huang, Yongxue Li, Mingdan Tan, Lan Wang, Peng Fang, Xi Chen, Nan Yan, and Hanjun Liu. An interpretable model based on graph learning for diagnosis of Parkinson's disease with voice-related EEG. *npj Digital Medicine*, 7(1), 2024. ISSN 2398-6352. doi: 10.1038/s41746-023-00983-9. URL http://dx.doi.org/10.1038/s41746-023-00983-9.

[30] Daniel A. Nation, Melanie D. Sweeney, Axel Montagne, Abhay P. Sagare, Lina M. D'Orazio, Maricarmen Pachicano, Farshid Sepehrband, Amy R. Nelson, David P. Buennagel, Michael G. Harrington, Tammie L. S. Benzinger, Anne M. Fagan, John M. Ringman, Lon S. Schneider, John C. Morris, Helena C. Chui, Meng Law, Arthur W. Toga, and Berislav V. Zlokovic. Blood–brain barrier breakdown is an early biomarker of human cognitive dysfunction. *Nature Medicine*, 25(2):270–276, 2019. ISSN 1546-170X. doi: 10.1038/s41591-018-0297-y. URL http://dx.doi.org/10.1038/s41591-018-0297-y.

[31] Melanie D. Sweeney, Abhay P. Sagare, and Berislav V. Zlokovic. Blood–brain barrier breakdown in Alzheimer disease and other neurodegenerative disorders. *Nature Reviews Neurology*, 14(3): 133–150, 2018. ISSN 1759-4766. doi: 10.1038/nrneurol.2017.188. URL http://dx.doi.org/10.1038/nrneurol.2017.188.

[32] Dan Z. Milikovsky, Jonathan Ofer, Vladimir V. Senatorov, Aaron R. Friedman, Ofer Prager, Liron Sheintuch, Netta Elazari, Ronel Veksler, Daniel Zelig, Itai Weissberg, Guy Bar-Klein, Evyatar Swissa, Erez Hanael, Gal Ben-Arie, Osnat Schefenbauer, Lyna Kamintsky, Rotem Saar-Ashkenazy, Ilan Shelef, Merav H. Shamir, Ilan Goldberg, Amir Glik, Felix Benninger, Daniela Kaufer, and Alon Friedman. Paroxysmal slow cortical activity in Alzheimer's disease and epilepsy is associated with blood-brain barrier dysfunction. *Science Translational Medicine*, 11(521), 2019. ISSN 1946-6242. doi: 10.1126/scitranslmed.aaw8954. URL http://dx.doi.org/10.1126/scitranslmed.aaw8954.

[33] Vladimir V. Senatorov, Aaron R. Friedman, Dan Z. Milikovsky, Jonathan Ofer, Rotem Saar-Ashkenazy, Adiel Charbash, Naznin Jahan, Gregory Chin, Eszter Mihaly, Jessica M. Lin, Harrison J. Ramsay, Ariana Moghbel, Marcela K. Preininger, Chelsy R. Eddings, Helen V. Harrison, Rishi Patel, Yishuo Shen, Hana Ghanim, Huanjie Sheng, Ronel Veksler, Peter H. Sudmant, Albert Becker, Barry Hart, Michael A. Rogawski, Andrew Dillin, Alon Friedman, and Daniela Kaufer. Blood-brain barrier dysfunction in aging induces hyperactivation of TGFβ signaling and chronic yet reversible neural dysfunction. *Science Translational Medicine*, 11 (521), 2019. ISSN 1946-6242. doi: 10.1126/scitranslmed.aaw8283. URL http://dx.doi.org/10.1126/scitranslmed.aaw8283.

[34] Dan Z. Milikovsky, Yotam Sharabi, Nir Giladi, Anat Mirelman, Ronen Sosnik, Firas Fahoum, and Inbal Maidan. Paroxysmal slow-wave events are uncommon in Parkinson's disease. *Sensors*, 23(2):918, 2023. ISSN 1424-8220. doi: 10.3390/s23020918. URL http://dx.doi.org/10.3390/s23020918.

[35] S. Amara Ganon, A. Friedman, and Y. Zigel. Electroencephalogram data-based analysis of paroxysmal slow wave events patterns in brain pathologies. In *2025 47th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, page 1–7. IEEE, 2025. doi: 10.1109/embc58623.2025.11254025. URL http://dx.doi.org/10.1109/embc58623.2025.11254025.

[36] Mengfan Li, Zhuoran Sun, Hairong Sun, Guochen Zhao, Bing Leng, Tengqun Shen, Song Xue, Huimin Hou, Zhenguang Li, and Jinbiao Zhang. Paroxysmal slow wave events are associated with cognitive impairment in patients with obstructive sleep apnea. *Alzheimer's Research & Therapy*, 14(1):200, 2022. doi: 10.1186/s13195-022-01153-x.

[37] Lindsey Power, Alon Friedman, and Timothy Bardouille. Atypical paroxysmal slow cortical activity in healthy adults: Relationship to age and cognitive performance. *Neurobiology of Aging*, 136:44–57, 2024. doi: 10.1016/j.neurobiolaging.2024.01.009.

[38] Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. In *5th International Conference on Learning Representations (ICLR)*, 2017. URL https://openreview.net/forum?id=SJU4ayYgl.

[39] Shahid Iqbal, Humaira Nisar, and Kim Ho Yeap. Quantitative EEG signatures of power and functional connectivity alterations in Alzheimer's disease and frontotemporal dementia.

Scientific Reports, 16(1), 2026. ISSN 2045-2322. doi: 10.1038/s41598-026-42452-9. URL http://dx.doi.org/10.1038/s41598-026-42452-9.

[40] Kassra Ghassemkhani, Kevin S. Saroka, and Blake T. Dotta. Evaluating EEG complexity and spectral signatures in Alzheimer's disease and frontotemporal dementia: evidence for rostrocaudal asymmetry. *npj Aging*, 11(1), 2025. ISSN 2731-6068. doi: 10.1038/s41514-025-00243-y. URL http://dx.doi.org/10.1038/s41514-025-00243-y.

[41] Mingkan Shen, Peng Wen, Bo Song, and Yan Li. Muti-band morlet mutual information functional connectivity for classifying Alzheimer's disease and frontotemporal dementia with a deep learning technique. *Computers in Biology and Medicine*, 197:111041, 2025. ISSN 0010-4825. doi: 10.1016/j.compbiomed.2025.111041. URL http://dx.doi.org/10.1016/j.compbiomed.2025.111041.

[42] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks. In *6th International Conference on Learning Representations (ICLR)*, 2018. URL https://openreview.net/forum?id=rJXMpikCZ.

[43] Mohamed Radwan, Pedro G. Lind, and Anis Yazidi. Frequency bands EEG biomarkers for dementia using graph neural networks. bioRxiv preprint (not peer reviewed), 2025. URL http://dx.doi.org/10.1101/2025.08.18.670945.

[44] Cornelis J. Stam, Guido Nolte, and Andreas Daffertshofer. Phase lag index: Assessment of functional connectivity from multi channel EEG and MEG with diminished bias from common sources. *Human Brain Mapping*, 28(11):1178–1193, 2007. ISSN 1097-0193. doi: 10.1002/hbm.20346. URL http://dx.doi.org/10.1002/hbm.20346.

[45] Martin Vinck, Robert Oostenveld, Marijn van Wingerden, Franscesco Battaglia, and Cyriel M.A. Pennartz. An improved index of phase-synchronization for electrophysiological data in the presence of volume-conduction, noise and sample-size bias. *NeuroImage*, 55(4):1548–1565, 2011. ISSN 1053-8119. doi: 10.1016/j.neuroimage.2011.01.055. URL http://dx.doi.org/10.1016/j.neuroimage.2011.01.055.

[46] Jiayi Cao, Bin Li, and Xiaoou Li. Identification of Alzheimer's disease brain networks based on EEG phase synchronization. *BioMedical Engineering OnLine*, 24(1), 2025. ISSN 1475-925X. doi: 10.1186/s12938-025-01361-0. URL http://dx.doi.org/10.1186/s12938-025-01361-0.

[47] Alexander Kamrud, Brett Borghetti, and Christine Schubert Kabban. The effects of individual differences, non-stationarity, and the importance of data partitioning decisions for training and testing of EEG cross-participant models. *Sensors*, 21(9):3225, 2021. doi: 10.3390/s21093225. URL https://doi.org/10.3390/s21093225.

[48] Russell A. Poldrack, Grace Huckins, and Gael Varoquaux. Establishment of best practices for evidence for prediction. *JAMA Psychiatry*, 77(5):534, 2020. doi: 10.1001/jamapsychiatry.2019.3671. URL https://doi.org/10.1001/jamapsychiatry.2019.3671.

[49] Gaël Varoquaux. Cross-validation failure: Small sample sizes lead to large error bars. *NeuroImage*, 180:68–77, 2018. doi: 10.1016/j.neuroimage.2017.06.061. URL https://doi.org/10.1016/j.neuroimage.2017.06.061.

[50] Alexandre Gramfort. MEG and EEG data analysis with MNE-Python. *Frontiers in Neuroscience*, 7, 2013. ISSN 1662-453X. doi: 10.3389/fnins.2013.00267. URL http://dx.doi.org/10.3389/fnins.2013.00267.

[51] Te-Won Lee, Mark Girolami, and Terrence J. Sejnowski. Independent component analysis using an extended infomax algorithm for mixed subgaussian and supergaussian sources. *Neural Computation*, 11(2):417–441, 1999. ISSN 1530-888X. doi: 10.1162/089976699300016719. URL http://dx.doi.org/10.1162/089976699300016719.

[52] Luca Pion-Tonachini, Ken Kreutz-Delgado, and Scott Makeig. ICLabel: An automated electroencephalographic independent component classifier, dataset, and website. *NeuroImage*, 198:181–197, 2019. ISSN 1053-8119. doi: 10.1016/j.neuroimage.2019.05.026. URL http://dx.doi.org/10.1016/j.neuroimage.2019.05.026.

[53] P. Welch. The use of fast Fourier transform for the estimation of power spectra: A method based on time averaging over short, modified periodograms. *IEEE Transactions on Audio and Electroacoustics*, 15(2):70–73, 1967. ISSN 0018-9278. doi: 10.1109/tau.1967.1161901. URL http://dx.doi.org/10.1109/TAU.1967.1161901.

[54] Gregory Lee, Ralf Gommers, Filip Waselewski, Kai Wohlfahrt, and Aaron O'Leary. PyWavelets: A python package for wavelet analysis. *Journal of Open Source Software*, 4(36):1237, 2019. ISSN 2475-9066. doi: 10.21105/joss.01237. URL http://dx.doi.org/10.21105/joss.01237.

[55] Bo Hjorth. EEG analysis based on time domain properties. *Electroencephalography and Clinical Neurophysiology*, 29(3):306–310, 1970. ISSN 0013-4694. doi: 10.1016/0013-4694(70)90143-4. URL http://dx.doi.org/10.1016/0013-4694(70)90143-4.

[56] Joshua S. Richman and J. Randall Moorman. Physiological time-series analysis using approximate entropy and sample entropy. *American Journal of Physiology-Heart and Circulatory Physiology*, 278(6):H2039–H2049, 2000. ISSN 1522-1539. doi: 10.1152/ajpheart.2000.278.6.h2039. URL http://dx.doi.org/10.1152/ajpheart.2000.278.6.H2039.

[57] Jukka-Pekka Onnela, Jari Saramäki, János Kertész, and Kimmo Kaski. Intensity and coherence of motifs in weighted complex networks. *Physical Review E*, 71(6), 2005. ISSN 1550-2376. doi: 10.1103/physreve.71.065103. URL http://dx.doi.org/10.1103/PhysRevE.71.065103.

[58] Vernon J Lawhern, Amelia J Solon, Nicholas R Waytowich, Stephen M Gordon, Chou P Hung, and Brent J Lance. Eegnet: a compact convolutional neural network for EEG-based brain–computer interfaces. *Journal of Neural Engineering*, 15(5):056013, 2018. ISSN 1741-2552. doi: 10.1088/1741-2552/aace8c. URL http://dx.doi.org/10.1088/1741-2552/aace8c.

[59] Arun Singh, Rachel Cole, Arturo Espinoza, Jim Cavanagh, and Nandakumar Narayanan. Rest eyes open (OpenNeuro accession ds004584, version 1.0.0), 2023. URL https://openneuro.org/datasets/ds004584/versions/1.0.0.

[60] Arun Singh, Rachel C. Cole, Arturo I. Espinoza, Jan R. Wessel, James F. Cavanagh, and Nandakumar S. Narayanan. Evoked mid-frontal activity predicts cognitive dysfunction in Parkinson's disease. *Journal of Neurology, Neurosurgery & Psychiatry*, 94(11):945–953, 2023. doi: 10.1136/jnnp-2022-330154. URL https://doi.org/10.1136/jnnp-2022-330154.

[61] Claude Nadeau and Yoshua Bengio. Inference for the generalization error. *Machine Learning*, 52(3):239–281, 2003. ISSN 1573-0565. doi: 10.1023/a:1024068626366. URL http://dx.doi.org/10.1023/A:1024068626366.

[62] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.

[63] Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7:1–30, 2006. URL https://jmlr.org/papers/v7/demsar06a.html.

[64] Yonglin Chen, Huabin Wang, Dailei Zhang, Liping Zhang, and Liang Tao. Multi-feature fusion learning for Alzheimer's disease prediction using EEG signals in resting state. *Frontiers in Neuroscience*, 17:1272834, 2023. doi: 10.3389/fnins.2023.1272834. URL https://doi.org/10.3389/fnins.2023.1272834.

[65] Jacob Kang, Hunseok Kang, and Jong-Hyeon Seo. CNN-based framework for Alzheimer's disease detection from EEG via dynamic mode decomposition. *Frontiers in Neuroinformatics*, 19:1706099, 2025. doi: 10.3389/fninf.2025.1706099. URL https://doi.org/10.3389/fninf.2025.1706099.

Supplementary material

Supplementary methods

Deep-learning baselines. To place the GCN against standard end-to-end EEG networks, EEGNet-8,2 [58] (temporal convolution of 64 taps, depthwise spatial convolution, separable convolution; 3331 parameters) and a compact 1D-CNN (three temporal convolution blocks over the 19 channels, global average pooling; 47 939 parameters) were trained on the same raw 10 s windows as P5, with the same 25 selection splits, optimiser, early-stopping rule, subject-equal weighting and subject-level aggregation as the GCN, so that every comparison is paired. EEGNet's max-norm weight constraints were replaced by the shared weight decay; no hyperparameter was tuned. Because the original EEGNet paper recommends dropout 0.25 for cross-subject classification, EEGNet was re-run once with that value (everything else unchanged; 0.23 GPU-hours) and compared with the original EEGNet and the GCN in its own Holm family, so the pre-specified comparisons above are unaffected.

The GCN learns its weights from the training subjects only and uses the validation subjects only to choose when to stop, whereas the classical models have no stopping rule. To compare models trained on the same subjects, we evaluate both under two regimes on the same 25 test splits: (i) *train only*, with the classical models fitted on the training subjects and the GCN as in the main grid; and (ii) *train + val*, with the classical models fitted on training and validation subjects and the GCN refitted on the same subjects for the number of epochs selected by early stopping in the corresponding main-grid run (525 refits, no validation monitoring).

Severity stratification. To ask whether the models only find advanced disease, patients were split at the conventional mild cut-point (MMSE ≥ 21 versus ≤ 20; MMSE is used for analysis only). For P3 $\times$ spatial on both fold sets we compared, within AD and within FTD, the probability assigned to the true class between strata (Mann–Whitney) and its rank correlation with MMSE (Spearman), with Holm correction over the 8 tests; the stratum sizes are those of the cohort and were not chosen.

Stopping-rule sensitivity. Early stopping on about 14 validation subjects often selected an early epoch. To test whether this shaped the conclusions, the full 7×3 grid was re-run on the 25 selection splits with patience 50 instead of 20, and with no early stopping at all (100 epochs on the training subjects, final weights kept; validation subjects unused), each paired with the main grid split by split.

Background-relative PSWE (exploratory). After observing that the fixed-threshold rate was strongly correlated with background slowing (Section 4.9), we added a post-hoc analysis, labelled as such throughout. Its definitions were fixed in the configuration before it was run on any real data. A relative PSWE is a run of at least 5 s in which a channel's MPF lies below that channel's own recording-median MPF by at least 2 Hz (primary), 1.5 Hz or 3 Hz, or by at least 2 or 3 robust standard deviations (1.4826$\times$ median absolute deviation).

Detector robustness. Two properties of the fixed-threshold detector could bias it: with 2 s windows the MPF is quantised to 0.5 Hz steps, so the 6 Hz threshold falls on a bin edge; and artefacts add low-frequency power. In a pre-specified robustness analysis we re-detected events with each segment zero-padded to an 8 s FFT (MPF on a 0.125 Hz grid, window unchanged), with amplitude-based artefact rejection (a channel's 2 s segment rejected if its peak-to-peak amplitude exceeded 200 µV or it was flat, SD < 0.5 µV; rejected seconds break events and are removed from the exposure time), and with both (primary), and repeated the statistics below. The unmodified

variant reproduced the main event counts exactly in both cohorts. The fixed-threshold statistics were also repeated on the recordings preprocessed without ICA (A8), and across a grid of event rules (A11: MPF thresholds of 5, 6 and 7 Hz crossed with minimum durations of 3, 5 and 8 s).

Collinearity of the slowing adjustment. Because PSWE rate and relative delta+theta power are strongly correlated, a model containing both can give unstable group coefficients. We therefore report the variance inflation of the slowing covariate given group, age and sex; refit the models on subjects whose slowing lies within the range observed in controls (common support); and compare groups on the ratio of observed to expected events, where the expectation comes from a negative-binomial model of events on slowing, age and sex fitted in controls only.

Compute. All experiments ran on free Kaggle sessions (4 vCPUs); the only GPU use was the two deep-learning baselines (0.34 GPU-hours of training on one Kaggle GPU, plus 0.23 for the EEGNet dropout sensitivity run). The 525-run main grid took 2.4 CPU-hours, and the 2,400 further runs (confirmation grid, ablations and permutations on the confirmation folds, and the train + val refits) took 9.0 CPU-hours. The stopping-rule sensitivity grid and the external cohort added 2,100 runs and 13.3 CPU-hours, and the learned-adjacency GCN 200 runs and 0.9 CPU-hours. Re-fitting the best cell for the occlusion analysis and for its reproducibility check took 50 further training runs (about 0.3 CPU-hours, estimated from the main-grid run times of that cell). The mixed models, equivalence and adjusted omnibus tests, severity analysis and calibration simulation (NB13–NB20) are statistical computations that took minutes each. Every run is logged with its configuration hash, code version and seed.

Supplementary results

By-subject mixed model (descriptive). In the by-subject mixed model (Table S2, Figure S1), representation had a strong effect on both fold sets (selection $F(6, 510) = 16.60$, $p < 0.001$; confirmation $F = 19.81$, $p < 0.001$). The marginal mean probability assigned to a subject's true class ranged from 0.373 (P2) and 0.391 (P4) to 0.525 (P6) and 0.529 (P3), and 13 of the 21 representation pairs differed after Holm correction on each fold set. Graph construction and its interaction with representation were also detectable at the subject level (selection: graph $F(2, 170) = 8.24$, $p < 0.001$; interaction $F(12, 1020) = 3.81$, $p < 0.001$; confirmation: $p = 0.002$ and $p < 0.001$) but negligible in size: the three graph types' marginal means spanned only 0.007 (0.461 spatial, 0.453 functional, 0.453 hybrid). Subjects accounted for 55% of the variance and subject × representation for 42%, while subject × graph variance was estimated at zero: individuals differ in which representation suits them, not in which graph does. A by-subject repeated-measures ANOVA on the same values gave essentially the same statistics (representation $F = 16.60$, graph $F = 8.22$; Greenhouse–Geisser $p < 0.001$ and $p < 0.001$). Modelling whether each subject was correctly classified, rather than the probability, weakened every effect but left the pattern intact: representation remained significant on both fold sets ($F = 4.02$, $p < 0.001$; confirmation $F = 5.19$, $p < 0.001$), with P3 again highest (mean proportion correct 0.598 and 0.611), whereas the graph effect was marginal on the selection folds ($p = 0.038$, range 0.013) and absent on the confirmation folds ($p = 0.917$).

Stopping-rule sensitivity. Early stopping selected a median of epoch 9. Re-running the whole grid with patience 50 (median epoch 17) or with no early stopping (100 epochs, final weights) left the conclusions unchanged (Table S10, Figure S5): P3 × spatial remained the best cell under every rule (0.603 and 0.593), the mean over all cells moved by at most 0.003, cell rankings were preserved (Spearman $\rho = 0.95$ and 0.81 against the main grid), representation stayed significant in

the uncorrected analysis, and no representation's score changed significantly under either rule (0 of 14 Holm-corrected Wilcoxon tests). Raw windows (P5), the configuration with most parameters, did not benefit from longer training (-0.005 and -0.016), and no GCN-versus-classical comparison became significant. Formal equivalence within ± 0.05 was shown for 3 of the 14 representation-level comparisons, all for patience 50; training without early stopping changed scores too variably for equivalence to be shown.

Severity. Mild patients were not detectably harder to classify. For P3 $\times$ spatial, mild AD (MMSE ≥ 21, $n = 8$) was classified correctly in 0.600 of runs on the selection folds and 0.475 on the confirmation folds, against 0.657 and 0.664 for moderate-to-severe AD ($n = 28$); FTD showed the opposite, small difference (0.400 mild, $n = 15$, against 0.300, $n = 8$; Table S7). None of the 8 tests was significant (smallest uncorrected $p = 0.207$), and within patients the true-class probability was unrelated to MMSE ($|\rho| \leq 0.06$). Mild AD was less often correct than moderate-to-severe AD for 7 of 7 representations, so a modest disadvantage for early AD is plausible, but with only 8 mild AD patients it cannot be estimated.

Background-relative PSWEs (exploratory). Defining PSWEs relative to each channel's own background reversed the group pattern. With the primary definition (a drop of at least 2 Hz below the channel median), AD showed *fewer* transient slowing events than controls (rate ratio 0.55, 0.35–0.89, $p = 0.014$), FTD a non-significant reduction (0.70), and the event rate was only weakly related to slowing ($\rho = -0.15$) and not to MMSE ($\rho = 0.04$). After additional adjustment for background MPF the AD effect was no longer significant (0.77, $p = 0.299$). The MAD-scaled variants gave the same direction more strongly and survived that adjustment (for 2 MAD: AD 0.43, $p = 0.014$; FTD 0.35, $p = 0.002$). Every variant pointed the same way: relative to their own background, patients with dementia had fewer, not more, transient slowing events.

Table S1: Status of every analysis. Planned: in the study plan committed before any data were analysed. Added: decided after the corresponding results were known, with definitions, variants and margins fixed in the configuration before the analysis was run. Exploratory: post hoc, labelled as such in the text.

Status	Analyses
Primary	Planned split-level ANOVA and pairwise tests on macro-F1, with the Nadeau–Bengio adjustment for overlapping training sets (adjustment added after review; Deviations)
Descriptive	By-subject mixed model and its cross-checks; unadjusted split-level ANOVA, Wilcoxon tests, Friedman/Nemenyi
Planned	Split-level ANOVA, pairwise Wilcoxon and Nadeau–Bengio tests, Friedman/Nemenyi; demographic, classical and deep-learning baselines; ablations A1–A12; leakage; permutation test; external validation; PSWE group, slowing and MMSE analyses (RQ6) and incremental value (RQ7)
Added	Confirmation folds; train + val regime; equivalence tests; Nadeau–Bengio-adjusted omnibus tests; graph-type equivalence within representations; stopping-rule sensitivity; PSWE detector robustness, rule grid and collinearity checks; learned adjacency; occlusion importance; mixed-model cross-checks (RM-ANOVA, correct/incorrect outcome); severity stratification; EEGNet with dropout 0.25; calibration simulations of the tests (one matched to the data); re-fit checks; out-of-bootstrap comparison of representations
Exploratory	Background-relative PSWE definitions

Table S2: By-subject linear mixed model (descriptive; it does not allow for overlapping training sets and is not calibrated for claims about new samples) of the probability assigned to each subject's true class: Wald F (numerator, denominator df) and p for each fixed-effect term, on the selection and confirmation folds of ds004504 and on the external cohort ds004584.

	Selection		Confirmation		ds004584	
Term	F (df)	p	F (df)	p	F (df)	p
Representation	16.60 (6, 510)	<0.001	19.81 (6, 510)	<0.001	20.03 (6, 882)	<0.001
Graph construction	8.24 (2, 170)	<0.001	6.37 (2, 170)	0.002	6.21 (2, 294)	0.002
Interaction	3.81 (12, 1020)	<0.001	4.13 (12, 1020)	<0.001	2.94 (12, 1764)	<0.001

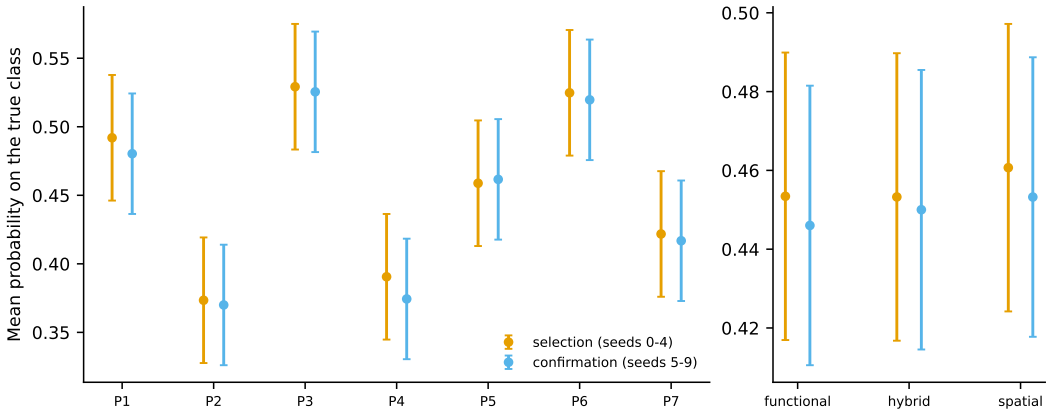


Figure S1: Model-based marginal mean probability on the true class (95% CI) for each representation (left, averaged over graph types) and each graph type (right), from the by-subject mixed model, on the selection and confirmation folds. The intervals include between-subject variation; the tests in Table S2 compare levels within subjects, which is why a difference much smaller than these intervals can still be detected.

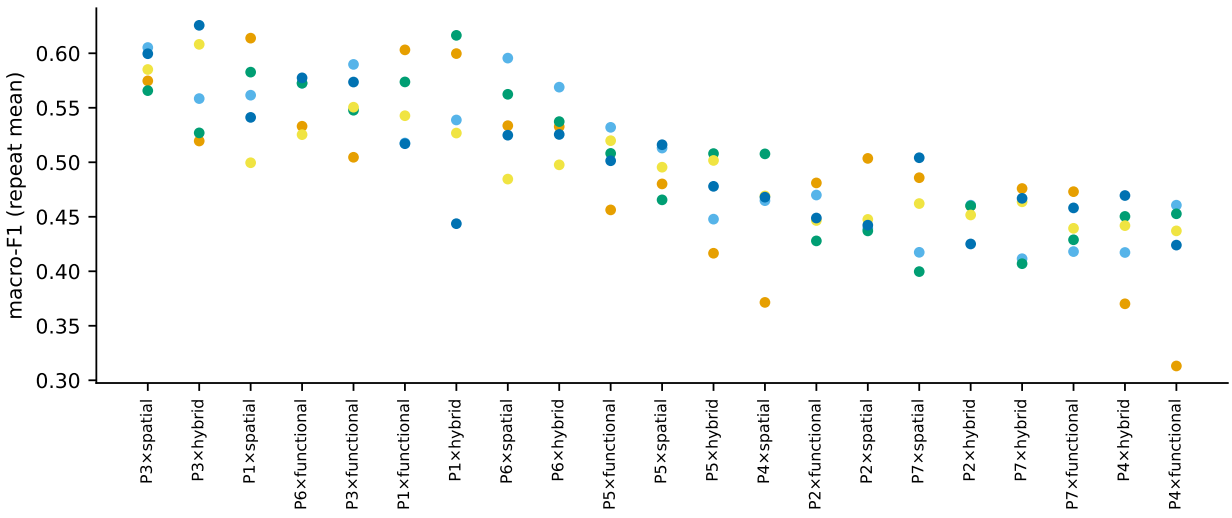


Figure S2: Mean macro-F1 of each repeat (seed) for every cell on the selection folds: variation due to the cross-validation partition and initialisation.

Table S3: Every representation's best GCN cell against its non-graph twins (same test splits, weighting and subject aggregation) under both training regimes, and the demographic floor. Macro-F1, mean ± SD over the 25 selection splits. Train only: models fitted on the training subjects (the GCN early-stopped on the validation subjects). Train + val: models fitted on training and validation subjects (the GCN refitted for its early-stopped epoch count). GCN graph type: best in each regime.

Representation	Model	Train only	Train + val
P1 spectral	GCN (spatial / functional)	0.560 ± 0.103	0.558 ± 0.110
	Logistic regression	0.536 ± 0.110	0.564 ± 0.096
	RBF-SVM	0.517 ± 0.126	0.560 ± 0.108
	Random forest	0.541 ± 0.116	0.588 ± 0.117
P2 connectivity	GCN (functional / spatial)	0.455 ± 0.108	0.453 ± 0.089
	Logistic regression	0.418 ± 0.083	0.417 ± 0.082
	RBF-SVM	0.459 ± 0.104	0.435 ± 0.089
	Random forest	0.437 ± 0.085	0.418 ± 0.055
P3 wavelet	GCN (spatial / spatial)	0.586 ± 0.112	0.613 ± 0.112
	Logistic regression	0.441 ± 0.132	0.466 ± 0.126
	RBF-SVM	0.513 ± 0.131	0.543 ± 0.122
	Random forest	0.536 ± 0.105	0.581 ± 0.092
P4 time domain	GCN (spatial / spatial)	0.456 ± 0.099	0.439 ± 0.097
	Logistic regression	0.441 ± 0.121	0.446 ± 0.112
	RBF-SVM	0.437 ± 0.109	0.483 ± 0.103
	Random forest	0.481 ± 0.119	0.487 ± 0.123
P5 raw	GCN (functional / spatial)	0.504 ± 0.087	0.484 ± 0.082
	EEGNet	0.530 ± 0.114	–
	EEGNet (dropout 0.25)	0.551 ± 0.116	–
	1D-CNN	0.523 ± 0.113	–
P6 fused	GCN (functional / spatial)	0.556 ± 0.099	0.580 ± 0.097
	Logistic regression	0.525 ± 0.129	0.514 ± 0.118
	RBF-SVM	0.561 ± 0.092	0.579 ± 0.098
	Random forest	0.540 ± 0.107	0.578 ± 0.118
P7 PSWE	GCN (spatial / spatial)	0.454 ± 0.082	0.456 ± 0.079
	Logistic regression	0.457 ± 0.108	0.456 ± 0.066
	RBF-SVM	0.485 ± 0.102	0.473 ± 0.088
	Random forest	0.498 ± 0.082	0.502 ± 0.077
Age + sex	Logistic regression	–	0.411 ± 0.120
	Stratified chance	–	0.302 ± 0.110

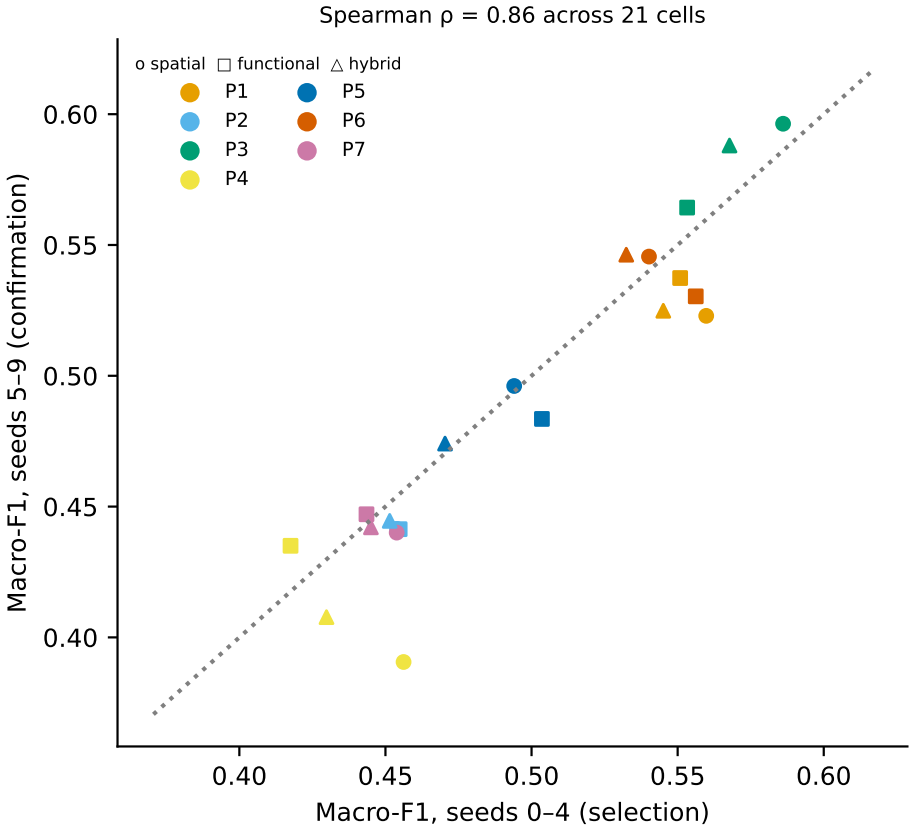


Figure S3: Every cell's mean macro-F1 on the selection folds (seeds 0–4) against the independent confirmation folds (seeds 5–9); dashed line, identity. Rank correlation: Spearman $\rho = 0.86$.

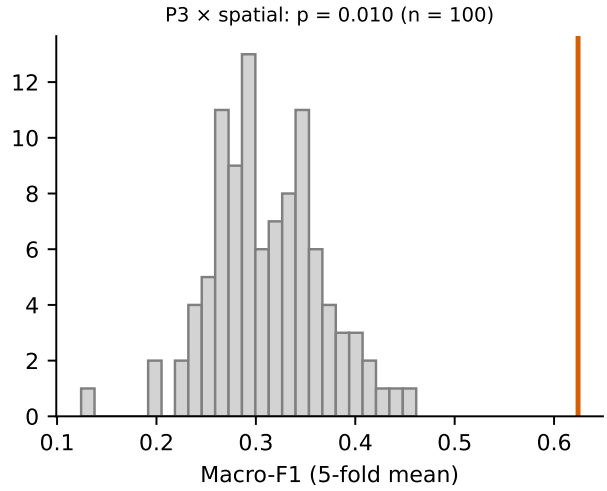


Figure S4: Label-permutation null distribution (100 permutations, first confirmation repeat's five folds) for the best cell; red line, observed macro-F1 on the same splits.

Table S4: Paired comparison of each representation's best GCN cell with its best non-graph twin under both training regimes (25 selection splits). Δ: GCN minus classical; p: Holm-corrected Wilcoxon and Nadeau–Bengio tests within each regime.

Regime	Representation	GCN	Classical	GCN F1	Classical F1	Δ	p_W	p_{NB}
train only	P1 spectral	GCN (spatial)	Random forest	0.560	0.541	+0.019	1.00	1.00
train only	P2 connectivity	GCN (functional)	RBF-SVM	0.455	0.459	-0.004	1.00	1.00
train only	P3 wavelet	GCN (spatial)	Random forest	0.586	0.536	+0.050	0.40	1.00
train only	P4 time domain	GCN (spatial)	Random forest	0.456	0.481	-0.025	1.00	1.00
train only	P6 fused	GCN (functional)	RBF-SVM	0.556	0.561	-0.004	1.00	1.00
train only	P7 PSWE	GCN (spatial)	Random forest	0.454	0.498	-0.044	0.31	1.00
train + val	P1 spectral	GCN (functional)	Random forest	0.558	0.588	-0.030	0.92	1.00
train + val	P2 connectivity	GCN (spatial)	RBF-SVM	0.453	0.435	+0.019	1.00	1.00
train + val	P3 wavelet	GCN (spatial)	Random forest	0.613	0.581	+0.033	0.80	1.00
train + val	P4 time domain	GCN (spatial)	Random forest	0.439	0.487	-0.049	0.38	1.00
train + val	P6 fused	GCN (spatial)	RBF-SVM	0.580	0.579	+0.000	1.00	1.00
train + val	P7 PSWE	GCN (spatial)	Random forest	0.456	0.502	-0.046	0.24	1.00

Table S5: Ablations on the selection folds (seeds 0–4), on which the two reference cells were chosen; layout as Table 4. These estimates are subject to selection bias and are given for comparison only.

	P3 × spatial		P3 × hybrid	
Ablation	Δ [95% CI]	p_{Holm}	Δ [95% CI]	p_{Holm}
Age/sex regressed out	-0.046 [-0.089, -0.003]	1.00	-0.044 [-0.093, +0.004]	1.00
No graph (identity)	-0.043 [-0.088, +0.003]	1.00	-0.024 [-0.064, +0.015]	1.00
Random graph (degree-preserving)	-0.037 [-0.075, +0.002]	1.00	-0.018 [-0.049, +0.013]	1.00
Fully connected	-0.066 [-0.113, -0.020]	1.00	-0.032 [-0.068, +0.005]	1.00
$k = 3$	-0.022 [-0.047, +0.004]	1.00	-0.009 [-0.026, +0.007]	1.00
$k = 6$	-0.009 [-0.033, +0.015]	1.00	+0.000 [-0.031, +0.032]	1.00
Binarised edges	-0.006 [-0.020, +0.008]	1.00	-0.023 [-0.057, +0.011]	1.00
Depth 1	-0.026 [-0.073, +0.020]	1.00	-0.006 [-0.042, +0.031]	1.00
Depth 3	-0.010 [-0.042, +0.022]	1.00	+0.007 [-0.018, +0.033]	1.00
GAT instead of GCN	-0.022 [-0.069, +0.025]	1.00	-0.021 [-0.044, +0.002]	1.00
No subject-equal weighting	+0.007 [-0.030, +0.044]	1.00	+0.014 [-0.016, +0.044]	1.00
+ P7 PSWE features	-0.024 [-0.059, +0.011]	1.00	+0.005 [-0.028, +0.038]	1.00
PSWE co-occurrence edges	-0.029 [-0.065, +0.007]	1.00	-0.011 [-0.044, +0.022]	1.00

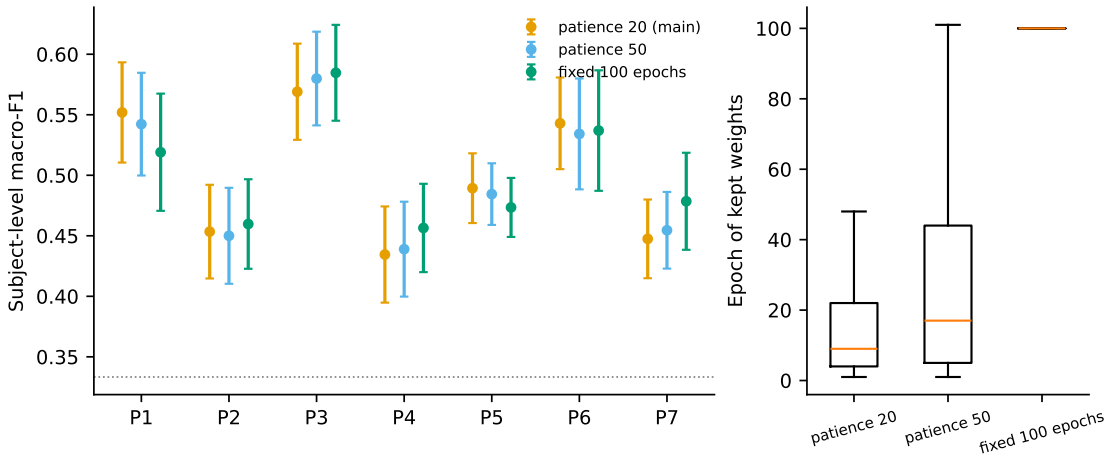


Figure S5: Stopping-rule sensitivity: macro-F1 per representation (mean and 95% CI over the 25 selection splits, averaged over graph types) under each stopping rule (left) and the epoch of the kept weights (right).

Table S6: Per-representation macro-F1 under each stopping rule (selection splits, averaged over graph types): change from the main grid and Holm-corrected Wilcoxon p (paired over the 25 splits).

	Patience 20	Patience 50			Fixed 100 epochs		
Representation	Macro-F1	Macro-F1	Δ	p_{Holm}	Macro-F1	Δ	p_{Holm}
P1	0.552	0.542	-0.010	0.56	0.519	-0.033	0.31
P2	0.453	0.450	-0.003	1.00	0.460	+0.006	1.00
P3	0.569	0.580	+0.011	0.97	0.585	+0.016	0.75
P4	0.435	0.439	+0.004	0.97	0.456	+0.022	0.35
P5	0.489	0.484	-0.005	1.00	0.473	-0.016	0.75
P6	0.543	0.534	-0.009	1.00	0.537	-0.006	1.00
P7	0.447	0.455	+0.007	1.00	0.479	+0.031	0.18

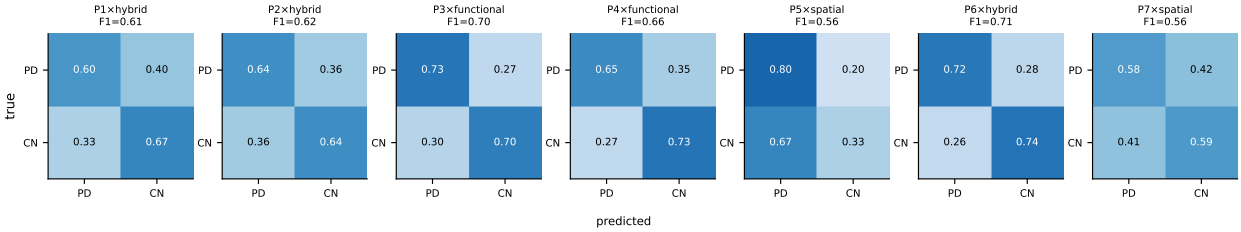


Figure S6: External cohort ds004584: row-normalised confusion matrices (summed over the 25 test splits) of the best graph type for each representation.

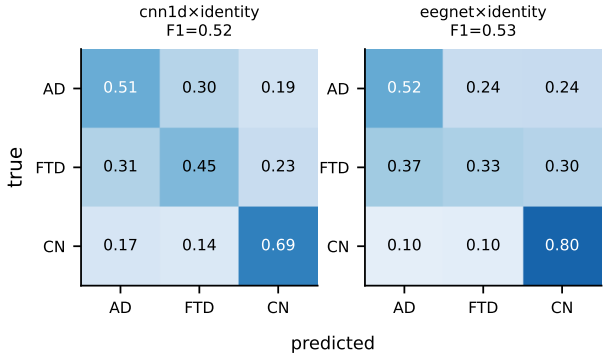


Figure S7: Deep-learning baselines on raw windows (ds004504, 25 selection splits): row-normalised confusion matrices of EEGNet and the 1D-CNN.

Table S7: Severity-stratified results for P3 × spatial (MMSE used for analysis only). Per stratum: number of patients, proportion of runs classified correctly, mean probability of the true class and proportion of runs predicted as controls, each averaged per patient over the five repeats. None of the Mann–Whitney or Spearman tests was significant after Holm correction.

		Selection folds				Confirmation folds			
Group	Stratum	n	Correct	p_{true}	Pred. CN	n	Correct	p_{true}	Pred. CN
AD	Mild (MMSE $\geq$ 21)	8	0.60	0.54	0.17	8	0.47	0.50	0.22
	Moderate–severe ($\leq$ 20)	28	0.66	0.62	0.19	28	0.66	0.61	0.18
FTD	Mild (MMSE $\geq$ 21)	15	0.40	0.36	0.24	15	0.45	0.37	0.29
	Moderate–severe ($\leq$ 20)	8	0.30	0.33	0.28	8	0.35	0.35	0.30

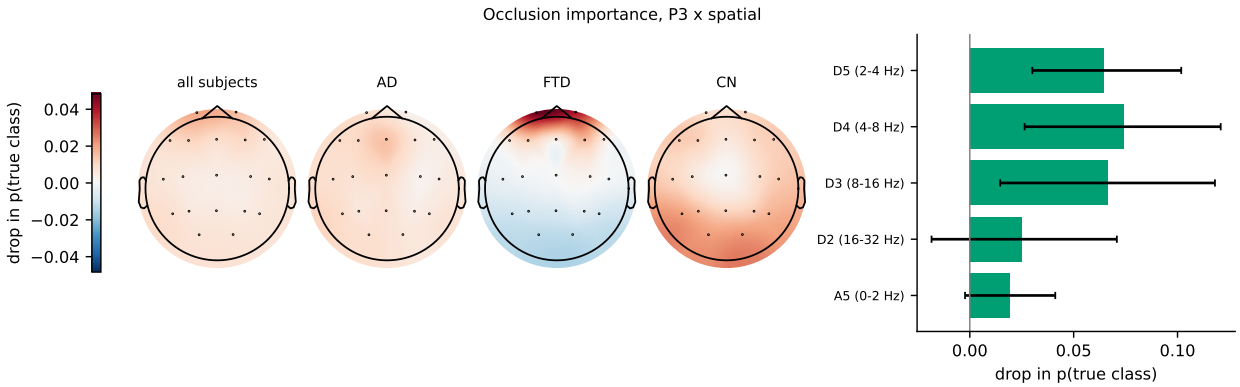


Figure S8: Occlusion importance for the best cell (P3 × spatial, 25 selection splits): drop in the probability assigned to the true class when one electrode (topomaps, overall and per class) or one wavelet sub-band on every electrode (bars, mean and subject-bootstrap 95% CI) is replaced by its training mean. Higher values mean the model relies more on that input.

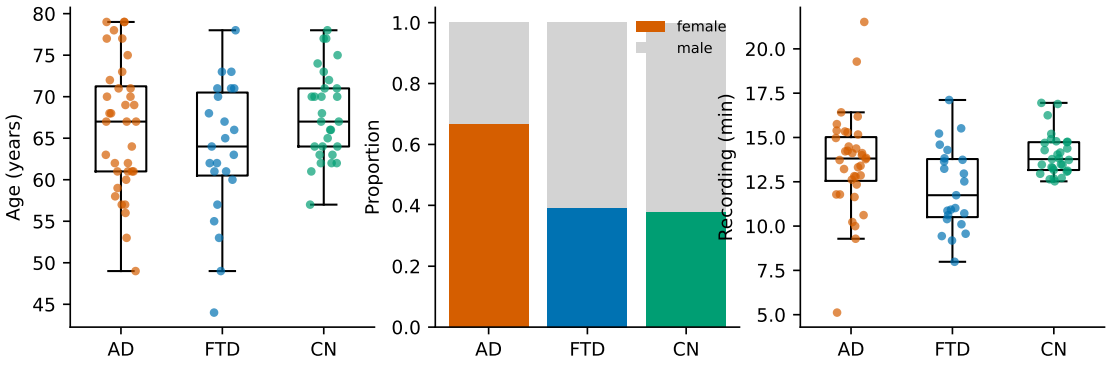


Figure S9: Cohort: age, sex and recording length by group.

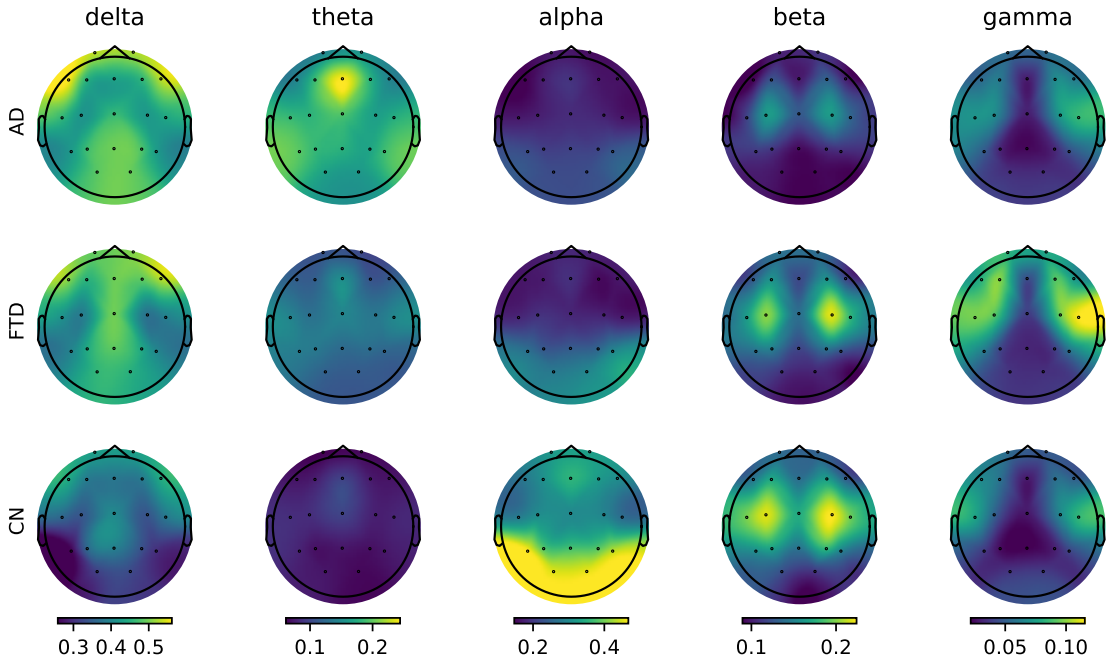


Figure S10: Group-mean relative band power per electrode (rows: groups; columns: bands; colour scales shared within each band).

Table S8: Subject-level performance of every cell (mean ± SD over the 25 selection splits, seeds 0–4). CI: subject-bootstrap 95% interval of macro-F1. BA: balanced accuracy; AUC: one-vs-rest macro ROC-AUC. Confirm.: macro-F1 of the same cell on the 25 independent confirmation splits (seeds 5–9).

Representation	Graph	Macro-F1	95% CI	BA	AUC	κ	Brier	Confirm.
P1 spectral	functional	0.551 ± 0.100	[0.507, 0.603]	0.565	0.724	0.363	0.578	0.537 ± 0.126
P1 spectral	hybrid	0.545 ± 0.117	[0.504, 0.600]	0.561	0.717	0.357	0.599	0.525 ± 0.114
P1 spectral	spatial	0.560 ± 0.103	[0.519, 0.616]	0.575	0.720	0.378	0.593	0.523 ± 0.120
P2 connectivity	functional	0.455 ± 0.108	[0.418, 0.507]	0.477	0.660	0.215	0.626	0.441 ± 0.077
P2 connectivity	hybrid	0.451 ± 0.090	[0.419, 0.507]	0.479	0.659	0.222	0.628	0.445 ± 0.089
P2 connectivity	spatial	0.454 ± 0.102	[0.420, 0.509]	0.483	0.653	0.225	0.629	0.441 ± 0.087
P3 wavelet	functional	0.553 ± 0.106	[0.503, 0.597]	0.560	0.724	0.359	0.598	0.564 ± 0.107
P3 wavelet	hybrid	0.568 ± 0.103	[0.522, 0.614]	0.578	0.735	0.384	0.582	0.588 ± 0.085
P3 wavelet	spatial	0.586 ± 0.112	[0.538, 0.628]	0.594	0.749	0.409	0.558	0.596 ± 0.089
P4 time domain	functional	0.418 ± 0.103	[0.382, 0.473]	0.436	0.623	0.163	0.668	0.435 ± 0.138
P4 time domain	hybrid	0.430 ± 0.104	[0.394, 0.488]	0.450	0.608	0.182	0.677	0.408 ± 0.145
P4 time domain	spatial	0.456 ± 0.099	[0.418, 0.511]	0.474	0.611	0.217	0.677	0.391 ± 0.152
P5 raw	functional	0.504 ± 0.087	[0.469, 0.556]	0.528	0.697	0.320	0.596	0.484 ± 0.120
P5 raw	hybrid	0.470 ± 0.086	[0.435, 0.514]	0.496	0.687	0.275	0.600	0.474 ± 0.089
P5 raw	spatial	0.494 ± 0.097	[0.458, 0.544]	0.526	0.706	0.318	0.594	0.496 ± 0.089
P6 fused	functional	0.556 ± 0.099	[0.506, 0.599]	0.571	0.731	0.370	0.605	0.530 ± 0.101
P6 fused	hybrid	0.532 ± 0.099	[0.490, 0.585]	0.554	0.724	0.342	0.617	0.546 ± 0.112
P6 fused	spatial	0.540 ± 0.094	[0.496, 0.589]	0.558	0.738	0.349	0.589	0.546 ± 0.111
P7 PSWE	functional	0.444 ± 0.085	[0.407, 0.498]	0.467	0.666	0.220	0.615	0.447 ± 0.109
P7 PSWE	hybrid	0.445 ± 0.084	[0.404, 0.495]	0.466	0.670	0.218	0.613	0.442 ± 0.109
P7 PSWE	spatial	0.454 ± 0.082	[0.415, 0.504]	0.472	0.662	0.225	0.620	0.440 ± 0.099

Table S9: Two-way repeated-measures ANOVA on subject-level macro-F1 (units: 25 test splits), on the selection folds and, independently, on the confirmation folds. ε: Greenhouse–Geisser; p: GG-corrected; η_p^2: partial eta squared.

		Selection (seeds 0–4)				Confirmation (seeds 5–9)			
Effect	df	F	ε	p	η_p^2	F	ε	p	η_p^2
Representation	6, 144	11.83	0.71	<0.001	0.33	13.57	0.77	<0.001	0.36
Graph construction	2, 48	4.03	0.96	0.026	0.14	0.03	0.91	0.957	0.00
Interaction	12, 288	1.43	0.60	0.194	0.06	1.58	0.56	0.149	0.06

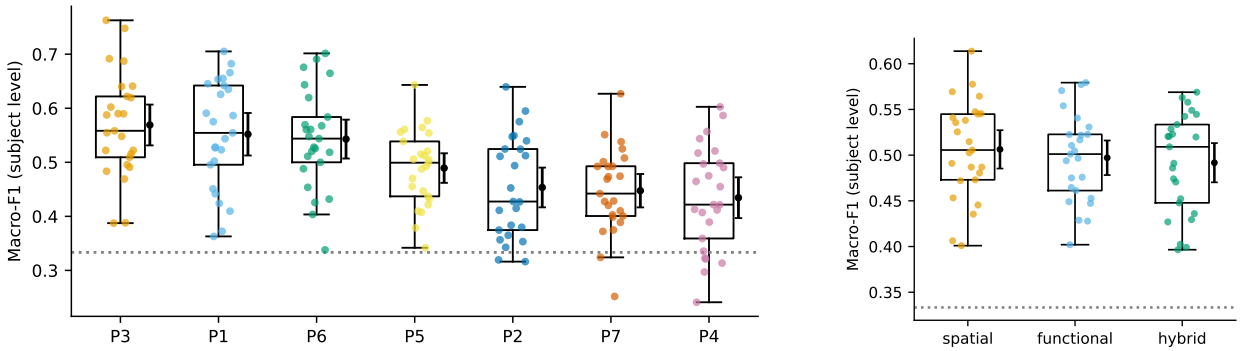


Figure S11: Subject-level macro-F1 by representation (left, averaged over graph types) and by graph construction (right, averaged over representations). Points are the 25 test splits; black markers, mean ± 95% CI; dotted line, 1/3.

Table S10: Stopping-rule sensitivity on the 25 selection splits: median epoch of the kept weights, best cell, macro-F1 of P3 × spatial, mean over the 21 cells, split-level ANOVA F for representation and graph, and Spearman correlation of the 21 cell means with the main grid.

Stopping rule	Epoch	Best cell	P3 × spatial	Mean	F_{rep}	F_{graph}	ρ
patience 20 (main)	9	P3 × spatial	0.586	0.498	11.83 (p <0.001)	4.03 (p =0.026)	1.00
patience 50	17	P3 × spatial	0.603	0.498	10.66 (p <0.001)	3.78 (p =0.032)	0.95
fixed 100 epochs	100	P3 × spatial	0.593	0.501	7.45 (p <0.001)	7.57 (p =0.002)	0.81

Table S11: Mean macro-F1 of each representation (averaged over graph types) and its rank on ds004504 (AD/FTD/CN, three classes) and on the external cohort ds004584 (PD/CN, two classes). Absolute values are not comparable across datasets; ranks are.

	ds004504		ds004584	
Representation	Macro-F1	Rank	Macro-F1	Rank
P1 spectral	0.552	2	0.594	5
P2 connectivity	0.453	5	0.614	4
P3 wavelet	0.569	1	0.689	2
P4 time domain	0.435	7	0.659	3
P5 raw	0.489	4	0.543	7
P6 fused	0.543	3	0.702	1
P7 PSWE	0.447	6	0.555	6

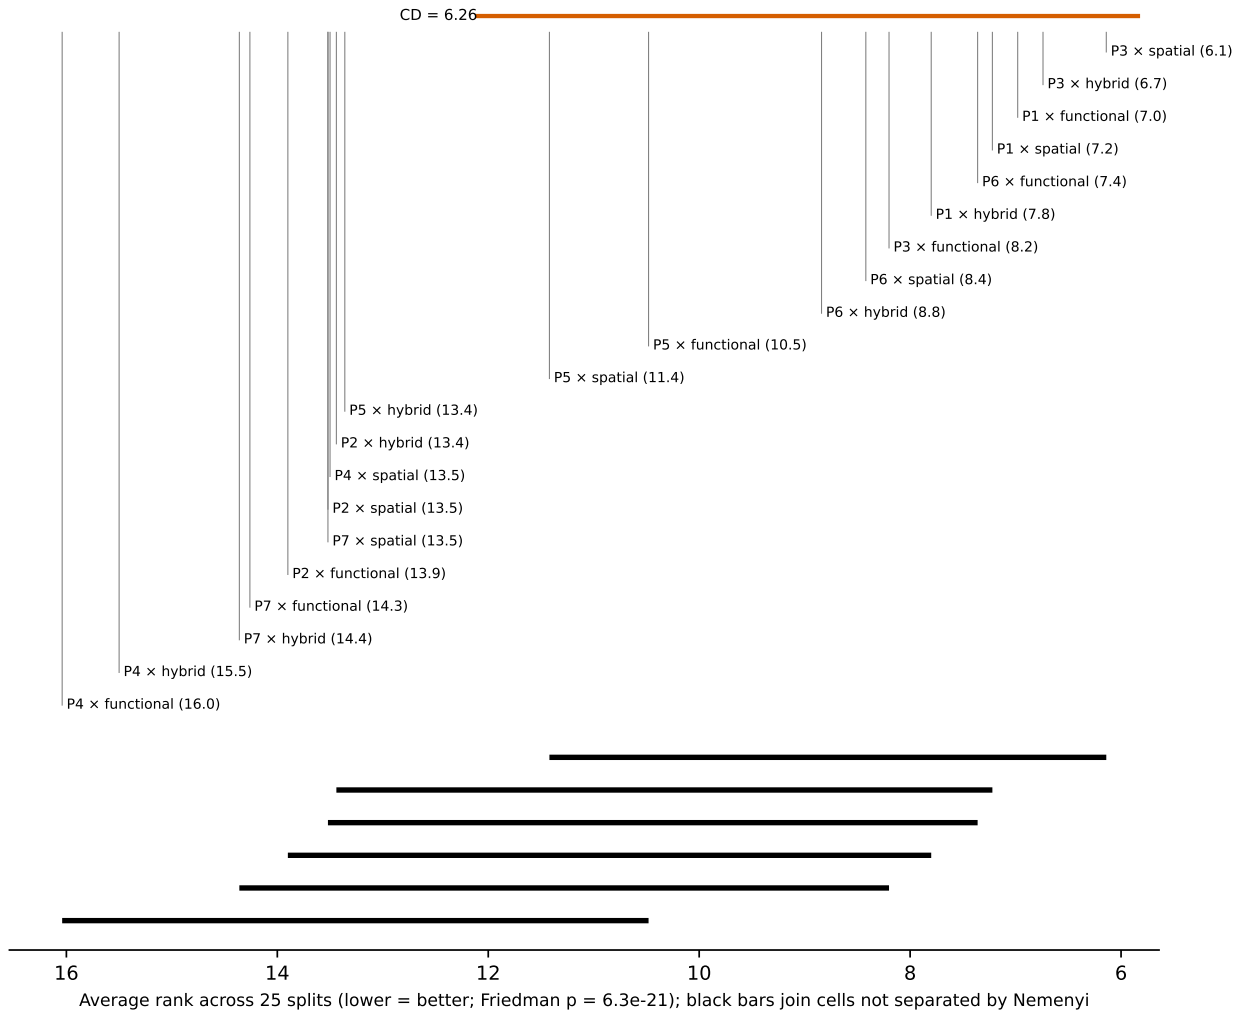


Figure S12: Critical-difference diagram over the 21 cells. Cells joined by a black bar are not separated by the Nemenyi test.

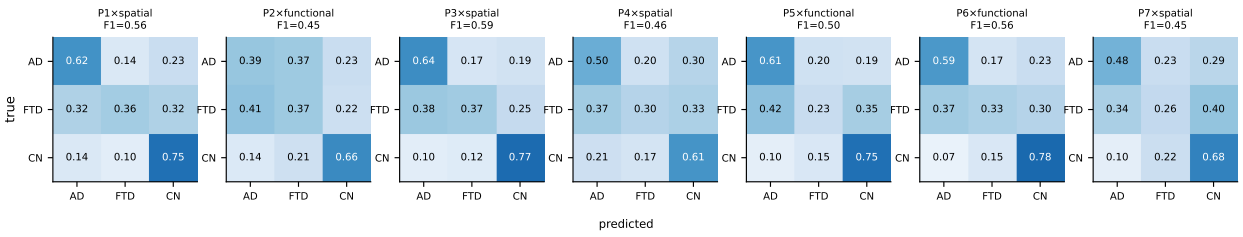


Figure S13: Row-normalised confusion matrices summed over the 25 test splits, for the best graph type of each representation.

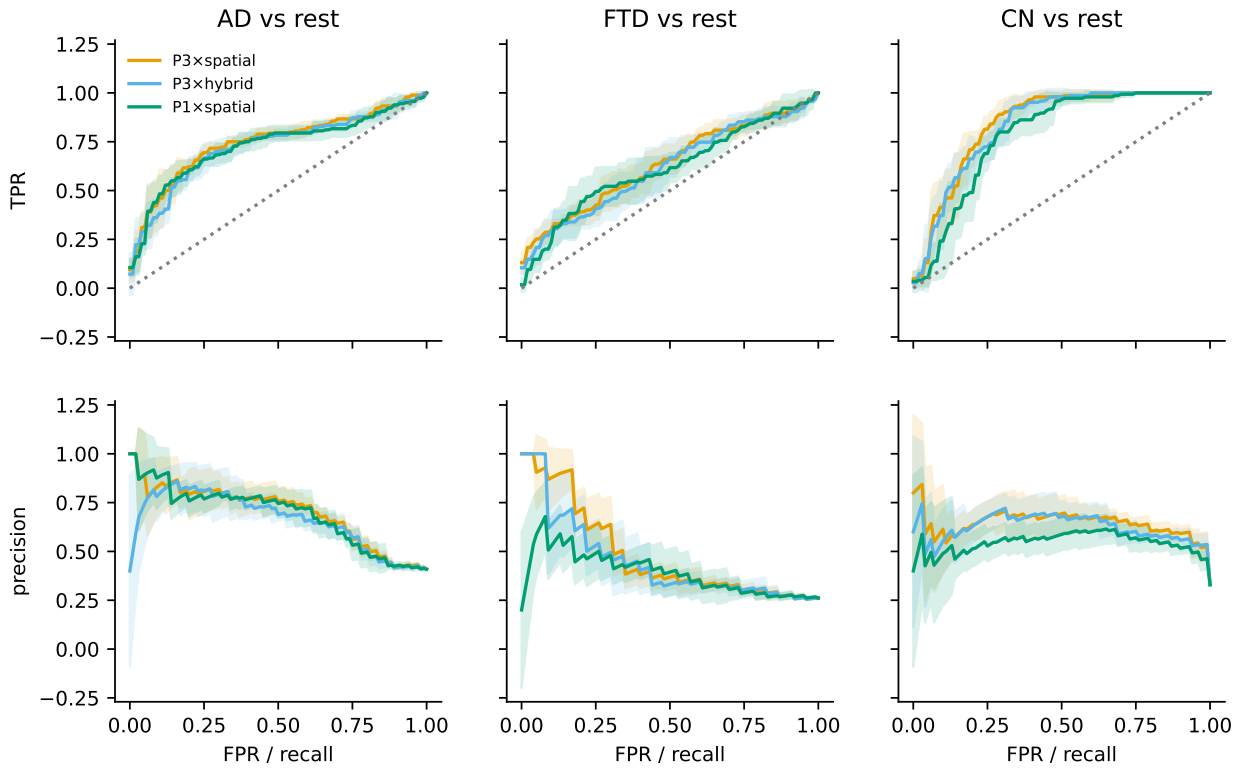


Figure S14: One-vs-rest ROC (top) and precision–recall (bottom) curves for the three best cells; mean ± SD across the five repeats.

Table S12: Equivalence tests for the comparisons read as "no difference": two one-sided tests against ±0.05 macro-F1 with the Nadeau–Bengio corrected standard error, Holm-corrected within each family. Max |Δ|: largest absolute mean difference; MDE: median difference detectable with 80% power.

Family	Folds	Comparisons	Equivalent	Max |Δ|	MDE
GCN vs non-graph twin	train only	6	0	0.050	0.19
GCN vs non-graph twin	train + val	6	0	0.049	0.18
Graph types (pairwise)	selection	3	3	0.015	0.04
Graph types (pairwise)	confirmation	3	3	0.002	0.05
Graph types within each representation	selection	21	0	0.039	0.11
Graph types within each representation	confirmation	21	0	0.044	0.12
Ablations	confirmation	32	0	0.046	0.14
Deep-learning baselines	selection	4	0	0.063	0.25
Stopping rules	selection	14	3	0.033	0.08
Learned vs fixed graph	selection	7	4	0.016	0.05
Learned vs fixed graph	confirmation	1	0	0.002	0.13
EEGNet dropout 0.25 vs 0.5	selection	1	0	0.021	0.20

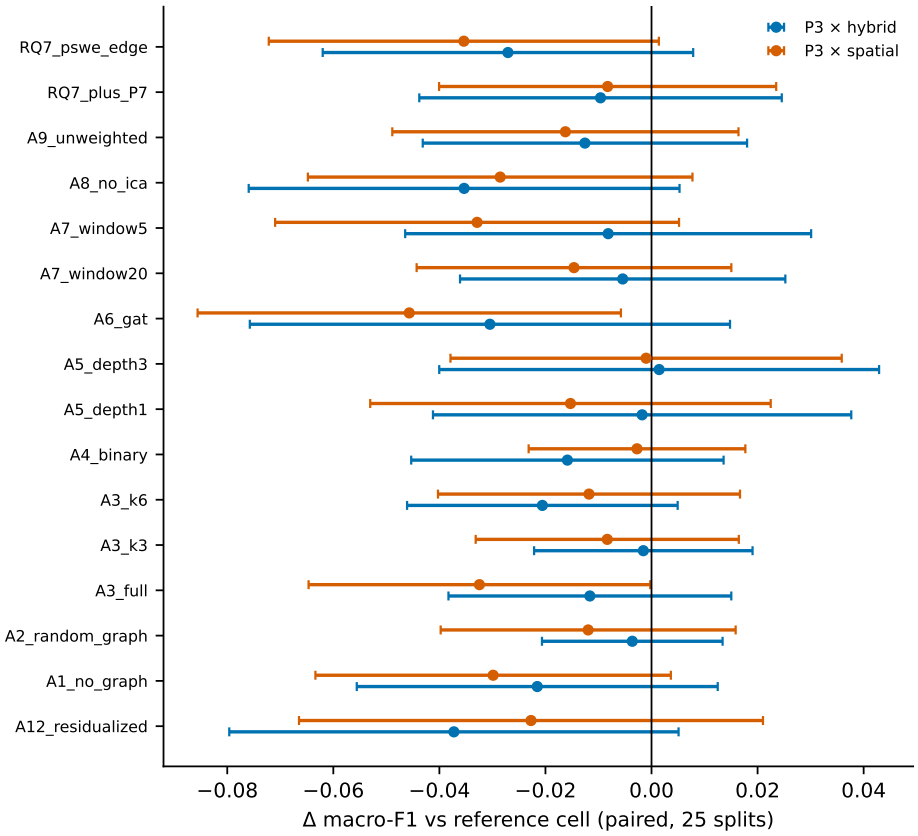


Figure S15: Ablations on the confirmation folds as paired differences from the two reference cells (mean and 95% CI over 25 splits). An asterisk would mark Nadeau–Bengio–Holm $p < 0.05$; none qualifies.

Table S13: PSWE event-rule grid (A11) on ds004504: age- and sex-adjusted rate ratios against controls without and with adjustment for relative delta+theta power, and the Spearman correlation of the PSWE rate with slowing, for each MPF threshold and minimum duration.

		Adjusted		+ slowing		
Threshold (Hz)	Min. duration (s)	AD	FTD	AD	FTD	ρ
5	3	2.25	1.79	0.56	0.77	0.88
5	5	2.76	1.78	0.41	0.65	0.82
5	8	3.21	1.37	0.24	0.35	0.80
6	3	2.29	1.87	0.66	0.79	0.91
6	5	2.96	2.07	0.50	0.67	0.88
6	8	3.75	1.99	0.30	0.40	0.86
7	3	2.14	1.84	0.78	0.88	0.89
7	5	2.95	2.29	0.64	0.75	0.90
7	8	3.95	2.75	0.40	0.49	0.91

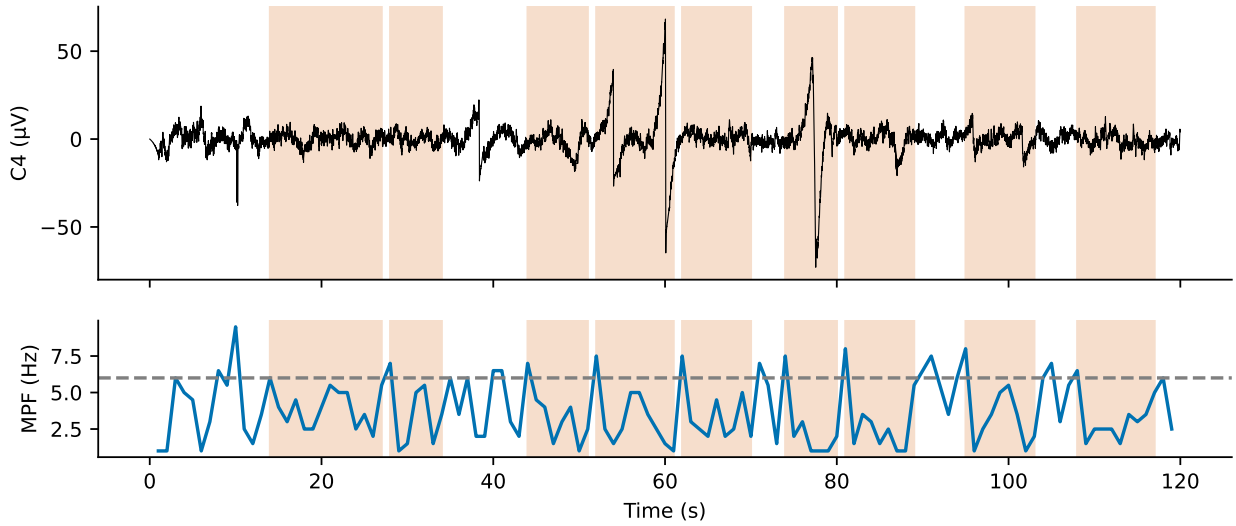


Figure S16: Fixed-threshold PSWE detection in the AD recording with the highest PSWE rate among the example shard (first 120 s, channel with most events). Top: EEG; bottom: median power frequency with the 6 Hz threshold; shaded, detected PSWEs. Because this subject's background is already slow, most of the recording is flagged — the pattern behind the collinearity with slowing.